



# CSCI-3753: Operating Systems

## Fall 2019

**Abigail Fernandes**

**Department of Computer Science**  
**University of Colorado Boulder**

# Programming Assignment 2



University of Colorado  
Boulder

# LKM – Quick Recap

What are the two main functions required in a LKM code?

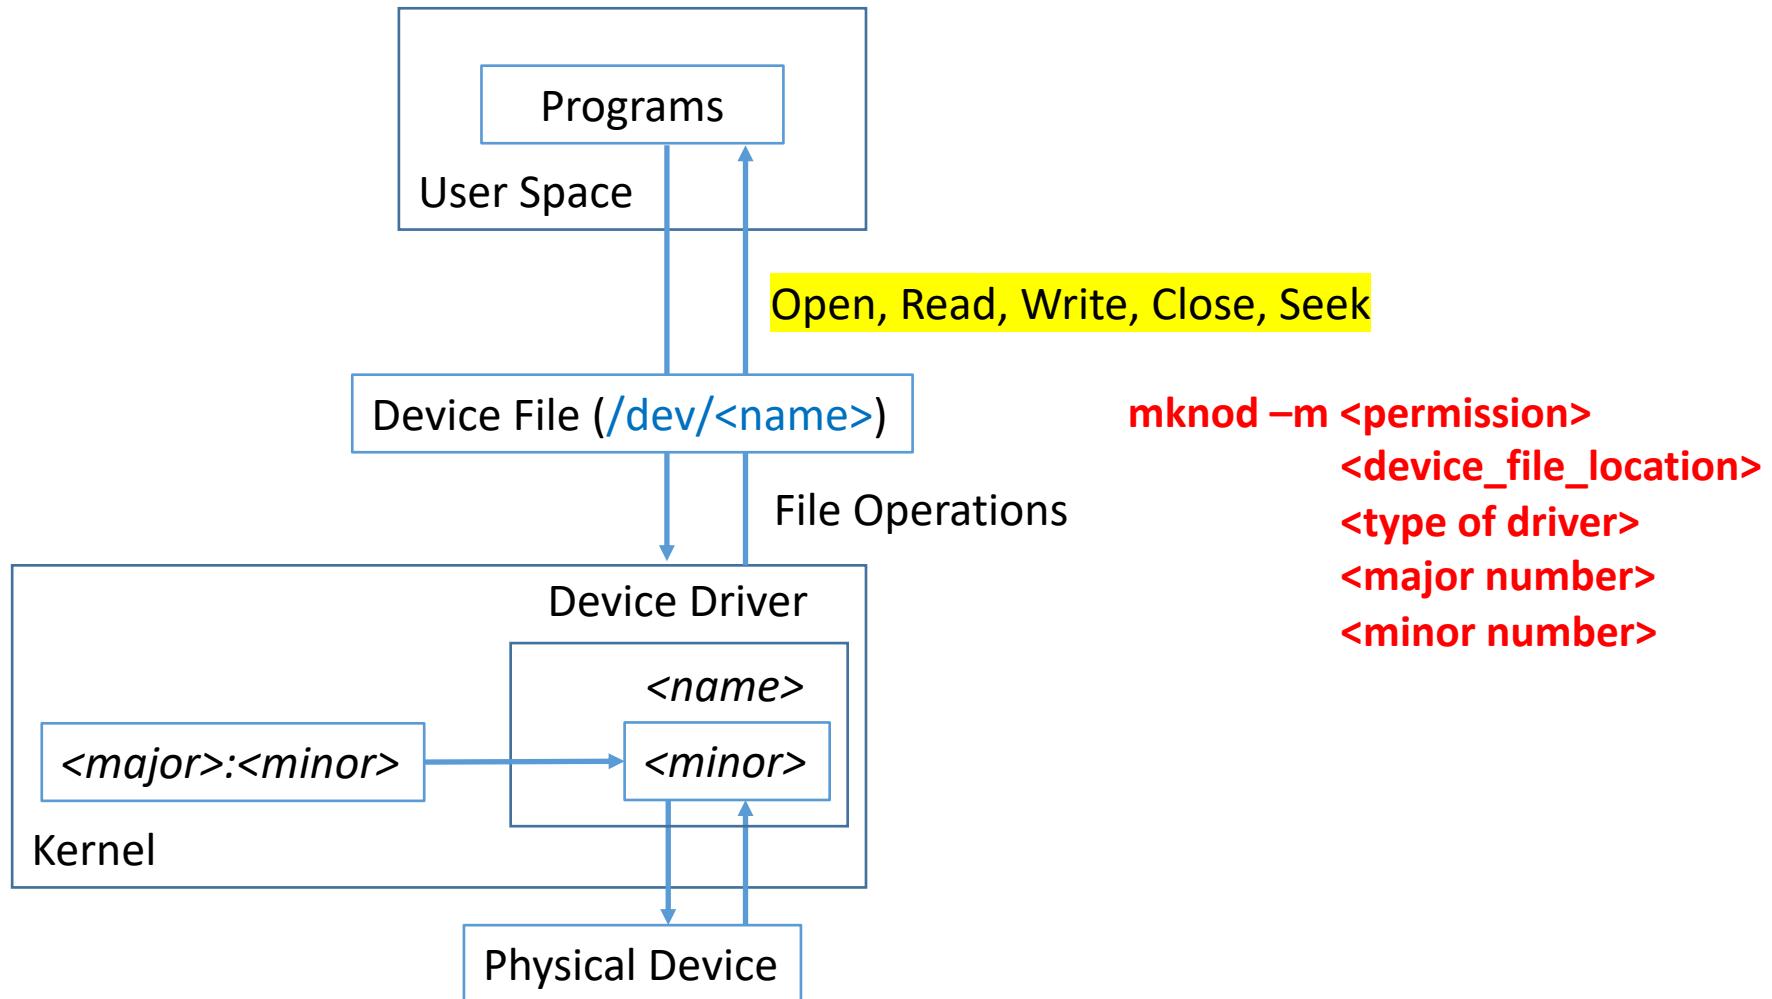
How do you compile the LKM code?

How do you load the module into the kernel?

How do you unload the module from the kernel?



# PA2 – Character Device Driver



# PA 2

1. Create a Device Driver Module (LKM)
2. Implement file operations (open, seek, read, write, release)
3. Make and load the module
4. Create a Device File for this Device
5. Test program that will allow you to perform different operations on the test file.



# File Operations Data Structure

- Defined in [/linux/fs.h](#)  
`vim /lib/modules/$(uname -r)/build/include/linux/fs.h`
- An open device is identified internally by a [file structure](#).
- The kernel uses the `file_operations` structure to access the driver's functions
- Each field in the structure must point to the function in the driver that implements a specific operation, or be left NULL for unsupported operations.



# File Operations Data Structure

```
struct file_operations {
    struct module *owner;
    loff_t (*llseek) (struct file *, loff_t, int);
    ssize_t (*read) (struct file *, char __user *, size_t, loff_t *);
    ssize_t (*write) (struct file *, const char __user *, size_t, loff_t *);
    ssize_t (*read_iter) (struct kiocb *, struct iov_iter *);
    ssize_t (*write_iter) (struct kiocb *, struct iov_iter *);
    int (*iterate) (struct file *, struct dir_context *);
    int (*iterate_shared) (struct file *, struct dir_context *);
    unsigned int (*poll) (struct file *, struct poll_table_struct *);
    long (*unlocked_ioctl) (struct file *, unsigned int, unsigned long);
    long (*compat_ioctl) (struct file *, unsigned int, unsigned long);
    int (*mmap) (struct file *, struct vm_area_struct *);
    unsigned long mmap_supported_flags;
    int (*open) (struct inode *, struct file *);
    int (*flush) (struct file *, fl_owner_t id);
    int (*release) (struct inode *, struct file *);
    int (*fsync) (struct file *, loff_t, loff_t, int datasync);
    int (*fasync) (int, struct file *, int);
    int (*lock) (struct file *, int, struct file_lock *);
    ssize_t (*sendpage) (struct file *, struct page *, int, size_t, loff_t *, int);
    unsigned long (*get_unmapped_area)(struct file *, unsigned long, unsigned long, unsigned long, unsigned long);
    int (*check_flags)(int);
    int (*setfl)(struct file *, unsigned long);
    int (*flock) (struct file *, int, struct file_lock *);
    ssize_t (*splice_write)(struct pipe_inode_info *, struct file *, loff_t *, size_t, unsigned int);
    ssize_t (*splice_read)(struct file *, loff_t *, struct pipe_inode_info *, size_t, unsigned int);
    int (*setlease)(struct file *, long, struct file_lock **, void **);
    long (*fallocate)(struct file *file, int mode, loff_t offset,
                      loff_t len);
    void (*show_fdinfo)(struct seq_file *m, struct file *f);
#endif CONFIG_MMU
    unsigned (*mmap_capabilities)(struct file *);
#endif
    ssize_t (*copy_file_range)(struct file *, loff_t, struct file *,
                             loff_t, size_t, unsigned int);
    int (*clone_file_range)(struct file *, loff_t, struct file *, loff_t,
                           u64);
    ssize_t (*dedupe_file_range)(struct file *, u64, u64, struct file *,
                                u64);
} __randomize_layout;
```



# System Device Major Numbers

<vim /home/kernel/linux-hwe-4.15.0/include/uapi/linux/major.h>

```
* SPDX-License-Identifier: GPL-2.0 WITH Linux-syscall-note */
#ifndef _LINUX_MAJOR_H
#define _LINUX_MAJOR_H

/*
 * This file has definitions for major device numbers.
 * For the device number assignments, see Documentation/admin-guide/devices.rst.
 */

#define UNNAMED_MAJOR      0
#define MEM_MAJOR          1
#define RAMDISK_MAJOR      1
#define FLOPPY_MAJOR        2
#define PTY_MASTER_MAJOR   2
#define IDE0_MAJOR         3
#define HD_MAJOR            IDE0_MAJOR
#define PTY_SLAVE_MAJOR    3
#define TTY_MAJOR           4
#define TTYAUX_MAJOR        5
#define LP_MAJOR            6
#define VCS_MAJOR           7
#define LOOP_MAJOR          7
#define SCSI_DISK0_MAJOR   8
#define SCSI_TAPE_MAJOR    9
#define MD_MAJOR            9
#define MISC_MAJOR          10
#define SCSI_CDROM_MAJOR   11
#define MUX_MAJOR            11 /* PA-RISC only */
#define XT_DISK_MAJOR       13
#define INPUT_MAJOR          13
#define SOUND_MAJOR          14
#define CDU31A_CDROM_MAJOR  15
#define JOYSTICK_MAJOR       15
#define GOLDSTAR_CDROM_MAJOR 16
#define OPTICS_CDROM_MAJOR  17
#define SANYO_CDROM_MAJOR   18
#define CYCLADES_MAJOR       19
#define CYCLADESAUX_MAJOR   20
```

# /dev Directory

```
$ ls -l /dev
brw-rw---- 1 root disk      8,   0 Dec 20 20:13 sda
crw-rw-rw- 1 root root     1,   3 Dec 20 20:13 null
srw-rw-rw- 1 root root      0 Dec 20 20:13 log
prw-r--r-- 1 root root      0 Dec 20 20:13 fdata
```

- Device Type
- Permissions
- Owner
- Group
- Major Device Number
- Minor Device Number
- Timestamp
- Device Name



# Documentation for devices

See Documentation/admin-guide/devices.txt -> Current devices and major numbers

vim /home/kernel/linux-hwe-4.15.0/Documentation/admin-guide/devices.txt

1 char	Memory devices	
	1 = /dev/mem	Physical memory access
	2 = /dev/kmem	Kernel virtual memory access
	3 = /dev/null	Null device
	4 = /dev/port	I/O port access
	5 = /dev/zero	Null byte source
	6 = /dev/core	OBSOLETE - replaced by /proc/kcore
	7 = /dev/full	Returns ENOSPC on write
	8 = /dev/random	Nondeterministic random number gen.
	9 = /dev/urandom	Faster, less secure random number gen.
	10 = /dev/aio	Asynchronous I/O notification interface
	11 = /dev/kmsg	Writes to this come out as printk's, reads export the buffered printk records.
	12 = /dev/oldmem	OBSOLETE - replaced by /proc/vmcore
1 block	RAM disk	
	0 = /dev/ram0	First RAM disk
	1 = /dev/ram1	Second RAM disk
	...	
	250 = /dev/initrd	Initial RAM disk



# Creating Device File for a Device Driver

```
sudo mknod -m <permission> <device_file_location> <type_of_driver>
<major_number> <minor_number>
```

Ex. sudo mknod -m 777 /dev/simple\_character\_device c 240 0

```
user@cu-cs-vm:/home/kernel/linux-hwe-4.15.0/Documentation/admin-guide$ sudo mknod -m 777 /dev/simple_character_device c 240 0
user@cu-cs-vm:/home/kernel/linux-hwe-4.15.0/Documentation/admin-guide$ ls /dev | grep simple
simple_character_device
user@cu-cs-vm:/home/kernel/linux-hwe-4.15.0/Documentation/admin-guide$ ls -l /dev | grep simple
crwxrwxrwx 1 root root 240, 0 Sep 19 22:59 simple_character_device
```



# PA2 – Requirements

1. Dynamically allocate kernel buffer to store the data written by the user
  - kmalloc()
    - Allocate memory for objects smaller than page size in the kernel at initialization time
  - kfree()
    - Free memory previously allocated using kmalloc() before exiting
2. Offset pointer in read/write function
3. Use copy\_to\_user and copy\_from\_user inside the device driver.



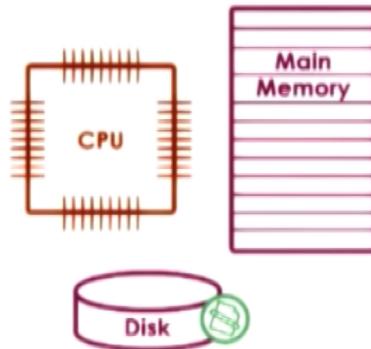
# Processes and Threads



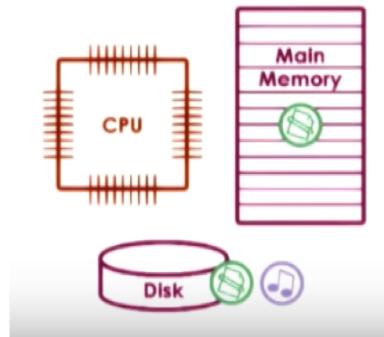
University of Colorado  
Boulder

# What is a process?

OS manages hardware on behalf of applications



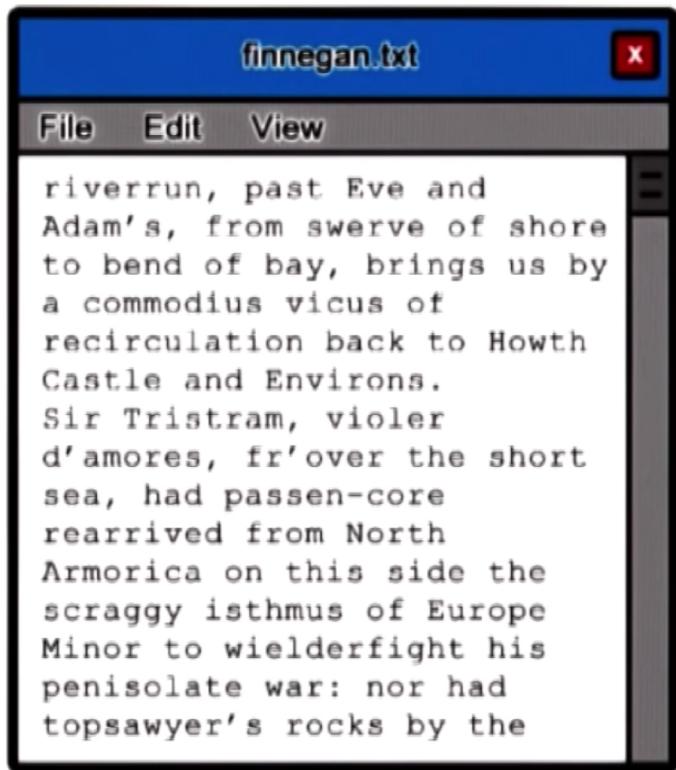
Application = program on disk, flash memory  
**(static entity)**



Process – state of a program when loaded in main memory  
**(active entity)**

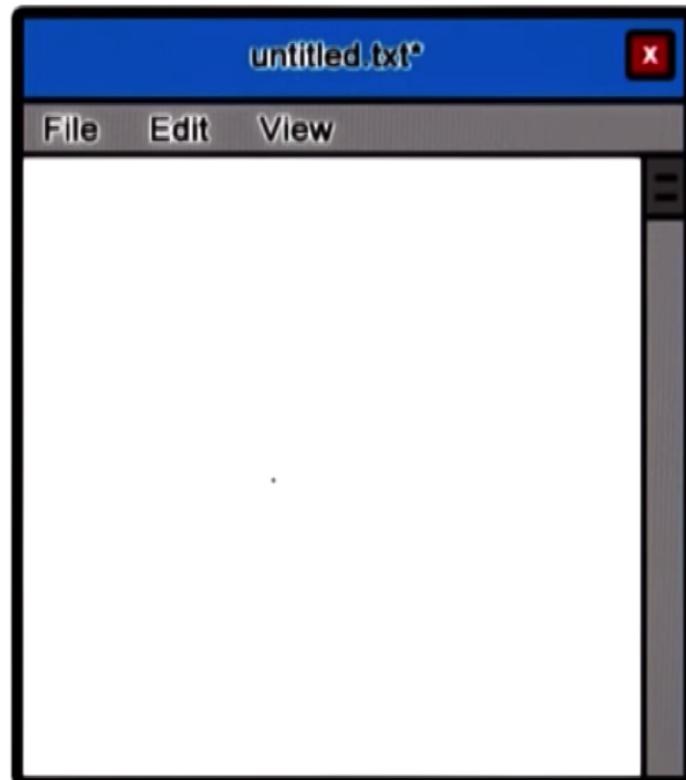


# Example of a process



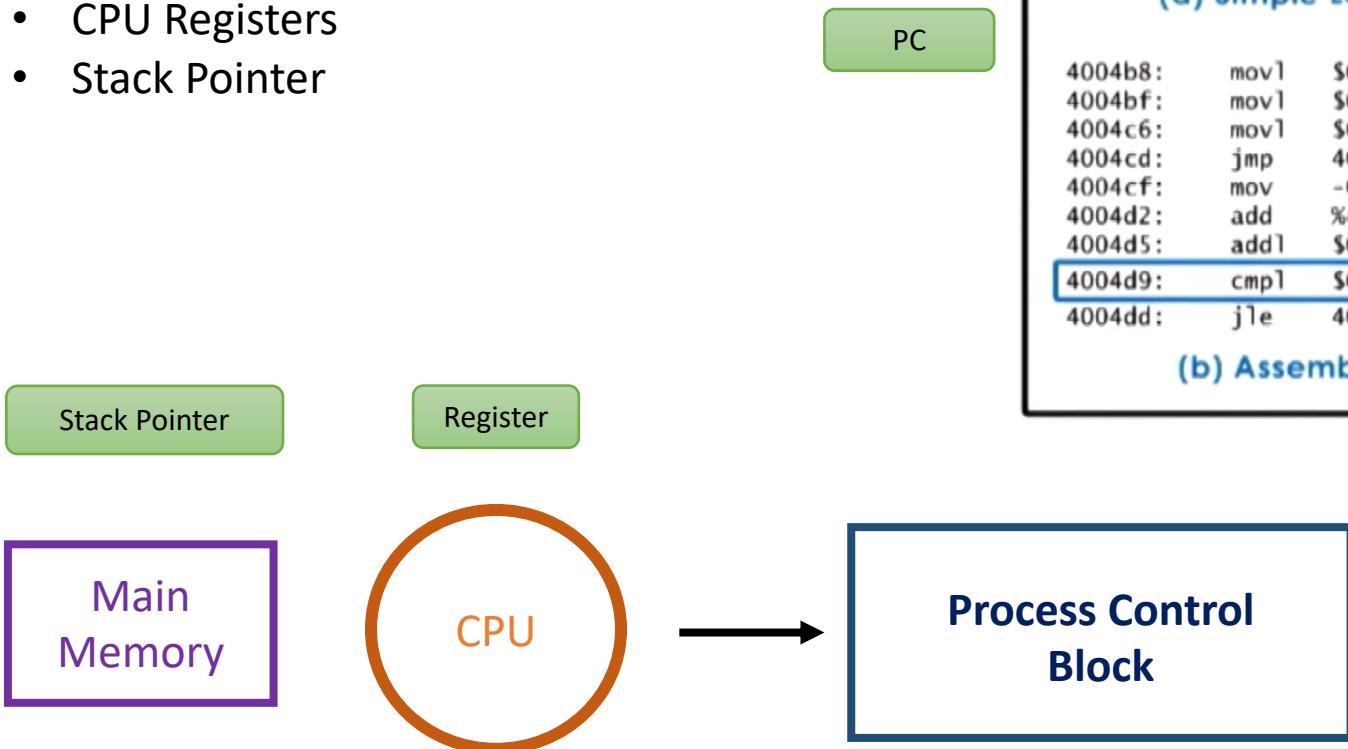
A screenshot of a text editor window titled "finnegan.txt". The window has a blue header bar with the title and a red close button. Below the header is a menu bar with "File", "Edit", and "View" options. The main content area contains the following text:

```
riverrun, past Eve and  
Adam's, from swerve of shore  
to bend of bay, brings us by  
a commodius vicus of  
recirculation back to Howth  
Castle and Environs.  
Sir Tristram, violer  
d'amores, fr'over the short  
sea, had passen-core  
rearrived from North  
Armorica on this side the  
scraggy isthmus of Europe  
Minor to wielderfight his  
penisolate war: nor had  
topsawyer's rocks by the
```



# How does the OS know what a process is doing?

- Program Counter
- CPU Registers
- Stack Pointer



```
sum = 0;
for (int i = 0; i < 10; ++i) {
    sum += i;
}
```

## (a) Simple Loop Code

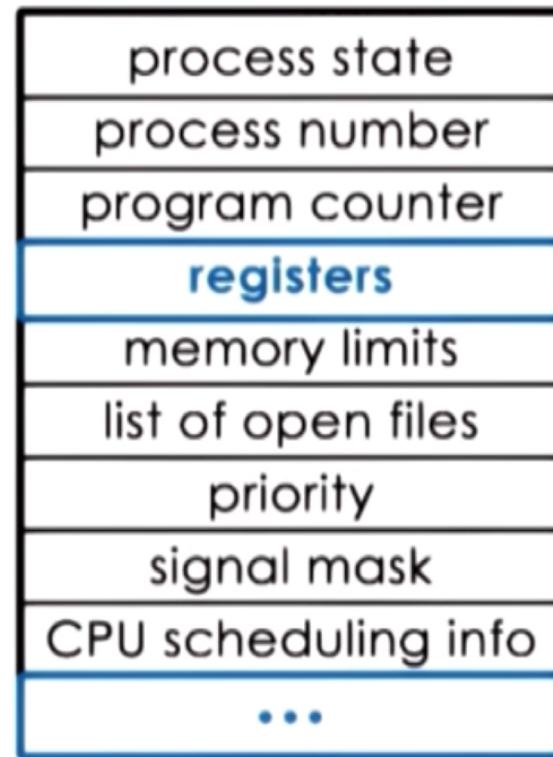
```
4004b8:    movl $0x0,-0x8(%rbp)
4004bf:    movl $0x0,-0x4(%rbp)
4004c6:    movl $0x0,-0x4(%rbp)
4004cd:    jmp  4004d9 <main+0x25>
4004cf:    mov  -0x4(%rbp),%cax
4004d2:    add  %cax,-0x8(%rbp)
4004d5:    addl $0x1,-0x4(%rbp)
4004d9:    cmpl $0x9,-0x4(%rbp)
4004dd:    jle   4004cf <main+0x1b>
```

## (b) Assembly Code

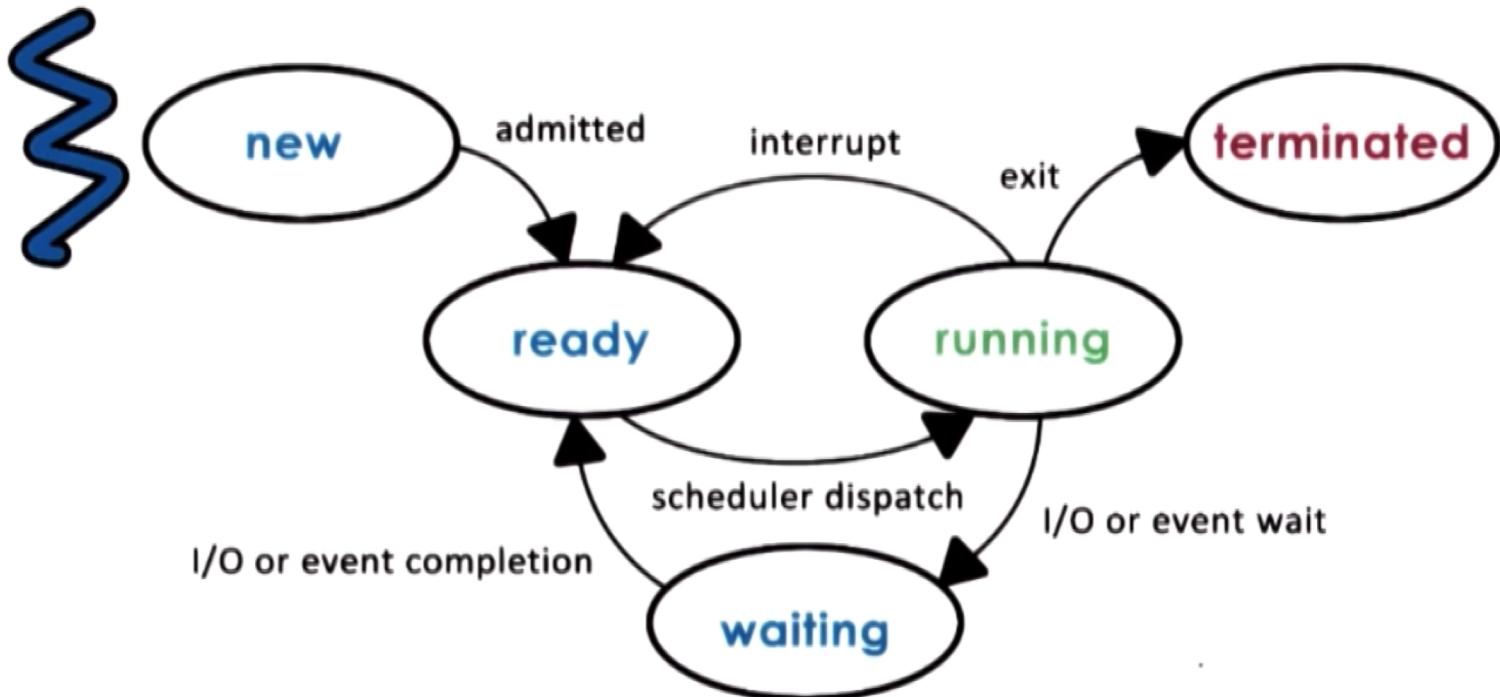


# Process Control Block

- PCB created when process is created
- Certain fields are updated when process state changes
- Other fields change very frequently



# Process Lifecycle: States



# Process Lifecycle: States

**New:** The process is being created.

**Running:** Instructions are being executed.

**Waiting:** The process is waiting for some event to occur (such as an I/O completion or reception of a signal).

**Ready:** The process is waiting to be assigned to a processor.

**Terminated:** The process has finished execution.



# Processes

Have an ID associated with it called the process ID (PID)

- **ps** command

- PID: Process ID
- TTY: Controlling terminal associated with the process
- STAT: Process status code
- TIME: Total CPU usage time
- CMD: Name of executable/command

```
$ ps
```

PID	TTY	STAT	TIME	CMD
41230	pts/4	Ss	00:00:00	bash
51224	pts/4	R+	00:00:00	ps



# /proc directory

```
user@cu-cs-vm:/home/kernel/linux-hwe-4.15.0/Documentation/admin-guide$ ls /proc
1      13    1407   1526   16718   177    19    202   219   25    30    42    843   957    fs      modules      thread-self
10     1334   1412   1540   17249   178    190   203   22    26    30805   43    844   96    interrupts  mounts      timer_list
1001   1339   1424   1545   17254   17853  1905   204   220   262   30836   6     85    971    iomem      mpt      tty
1019   1341   1427   1552   17255   17858  19059   205   221   263   31    7     852    acpi     iports      mtrr      uptime
1024   1350   1443   1559   17258   179    191   206   222   264   32    70946   855    buddyinfo  irq      net      version
1025   1359   1449   1573   17259   18     192   207   223   26482  32694   72513   857    bus      kallsyms  pagetypeinfo  version_signature
1028   1370   1453   1575   17262   180    1924   208   224   265   327    72684   86    cgroups  kcore      partitions  vmallocinfo
105     1374   1457   1580   173     181    193   209   225   267   33    72917   87    cmdline  keys      sched_debug  vmstat
1050   1384   1458   1581   17337   182    194   21    226   268   34    72926   878    consoles  key-users  schedstat  zoneinfo
11     1388   1467   1584   17370   183    195   210   227   27    35    73003   88    cpuinfo  kmsg      scsi
1171   1390   1471   1593   174     184    196   211   228   27059  36    73096   880    crypto   kpagegroup self
12     1392   1476   16     17414   185    197   212   229   27061  364    73163   886    devices  kpagecount slabinfo
122     1395   15    1664   17431   186    198   213   230   2743   369    73174   89    diskstats  kpageflags softirqs
1222   1399   1504   16701   175    187    199   214   231   28    37    73187   9     dma      loadavg  stat
1278   14     1517   16710   176    188    2     215   232   28518  393    742    90    driver   locks      swaps
1284   1400   1520   16715   17682  18882  20     216   233   29    4     8     906    execdomains mdstat      sys
1295   1402   1522   16716   17683  18886  200    217   234   290   405    829    907    fb      meminfo  sysrq-trigger
1297   1406   1523   16717   17686  189    201   218   24    291   41    838    929    filesystems misc      sysvipc
```

```
File Edit view terminal Tabs Help
user@cu-cs-vm:/home/kernel/linux-hwe-4.15.0/Documentation/admin-guide$ sudo ls /proc/1
attr      cmdline      environ      io      mem      ns      pagemap      sched      smaps_rollup      syscall      wchan
autogroup  comm      exe      limits      mountinfo      numa_maps      patch_state      schedstat      stack      task
auxv      coredump_filter      fd      loginuid      mounts      oom_adj      personality      sessionid      stat      timers
cgroup      cpuset      fdinfo      map_files      mountstats      oom_score      projid_map      setgroups      statm      timerslack_ns
clear_refs  cwd      gid_map      maps      net      oom_score_adj      root      smaps      status      uid_map
```



## /proc directory

- Contain virtual files and numbered directories corresponding to each running process
- Directory names = process ids
- When a process ends, its directory in /proc disappears automatically.
- Files in a numbered directory
  - cmdline
  - cwd
  - environ
  - fd
  - maps, statm, mem; stat, status



# /proc directory

Some typical virtual files that provide  
Hardware information

- [/proc/cpuinfo](#): identifies the type of processor used by your system
- [/proc/iomem](#): shows you the current map of the system's memory for each physical device
- [/proc/meminfo](#)
- [/proc/interrupts](#)

File-related info

- [/proc/filesystems](#): displays a list of the file system types currently supported by the kernel
- [/proc/partitions](#)

Kernel configuration parameters

- [/proc/cmdline](#): shows the parameters passed to the kernel at the time it is started
- [/proc/sys](#)



# top command

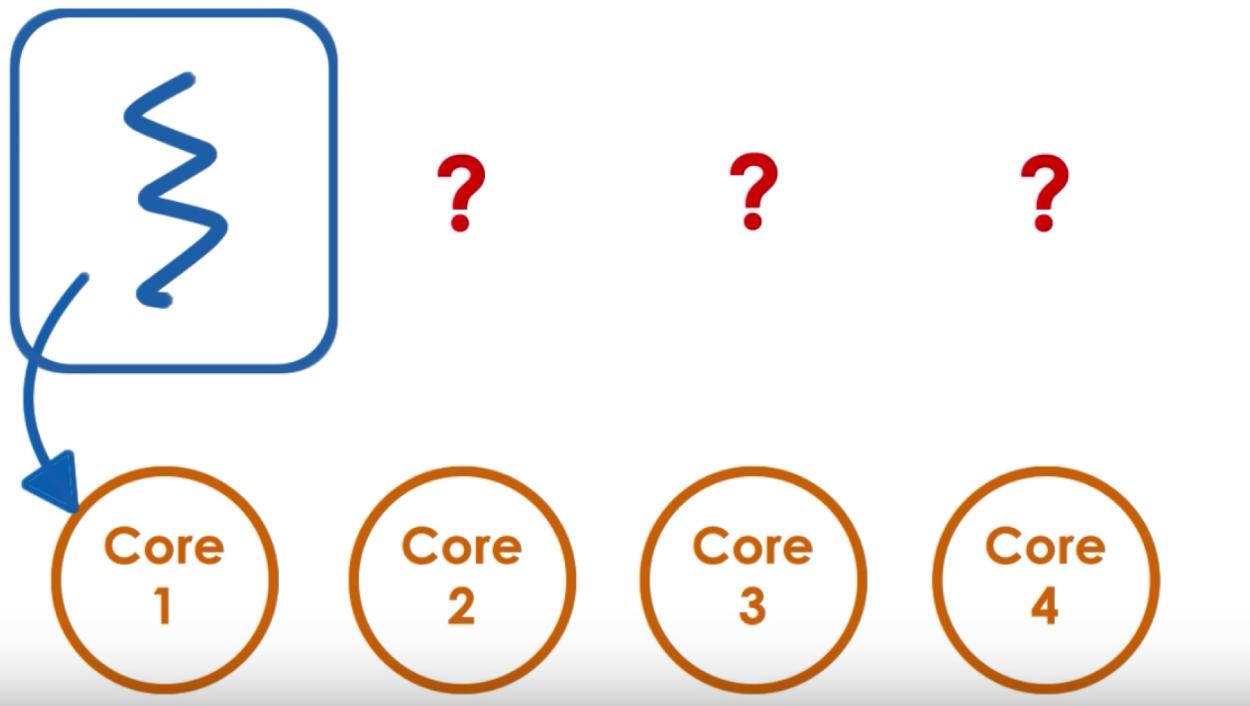
- Display a list of processes with their resource usage.

```
Terminal - user@cu-cs-vm:~  
File Edit View Terminal Tabs Help  
top - 01:10:46 up 36 min, 2 users, load average: 0.00, 0.00, 0.00  
Tasks: 208 total, 1 running, 142 sleeping, 0 stopped, 0 zombie  
%Cpu(s): 0.2 us, 0.2 sy, 0.0 ni, 99.7 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st  
KiB Mem : 2017320 total, 648132 free, 362208 used, 1006980 buff/cache  
KiB Swap: 998396 total, 998396 free, 0 used. 1419228 avail Mem  
  
 PID USER      PR  NI    VIRT    RES    SHR S %CPU %MEM     TIME+ COMMAND  
1362 user      20   0 439920  35304  28776 S  0.3  1.8  0:02.20 vmtoolsd  
2333 user      20   0 49288  3736  3096 R  0.3  0.2  0:00.09 top  
  1 root       20   0 119996  6136  4020 S  0.0  0.3  0:03.12 systemd  
  2 root       20   0      0      0      0 S  0.0  0.0  0:00.01 kthreadd  
  3 root       20   0      0      0      0 I  0.0  0.0  0:00.06 kworker/0+  
  4 root      0 -20      0      0      0 I  0.0  0.0  0:00.00 kworker/0+  
  6 root      0 -20      0      0      0 I  0.0  0.0  0:00.00 mm_percpu+  
  7 root       20   0      0      0      0 S  0.0  0.0  0:00.02 ksoftirqd+  
  8 root       20   0      0      0      0 I  0.0  0.0  0:00.09 rcu_sched  
  9 root       20   0      0      0      0 I  0.0  0.0  0:00.00 rcu_bh  
 10 root      rt  0      0      0      0 S  0.0  0.0  0:00.00 migration+  
 11 root      rt  0      0      0      0 S  0.0  0.0  0:00.00 watchdog/0  
 12 root      20   0      0      0      0 S  0.0  0.0  0:00.00 cpuhp/0  
 13 root      20   0      0      0      0 S  0.0  0.0  0:00.00 cpuhp/1  
 14 root      rt  0      0      0      0 S  0.0  0.0  0:00.00 watchdog/1  
 15 root      rt  0      0      0      0 S  0.0  0.0  0:00.00 migration+  
 16 root      20   0      0      0      0 S  0.0  0.0  0:00.05 ksoftirqd+
```

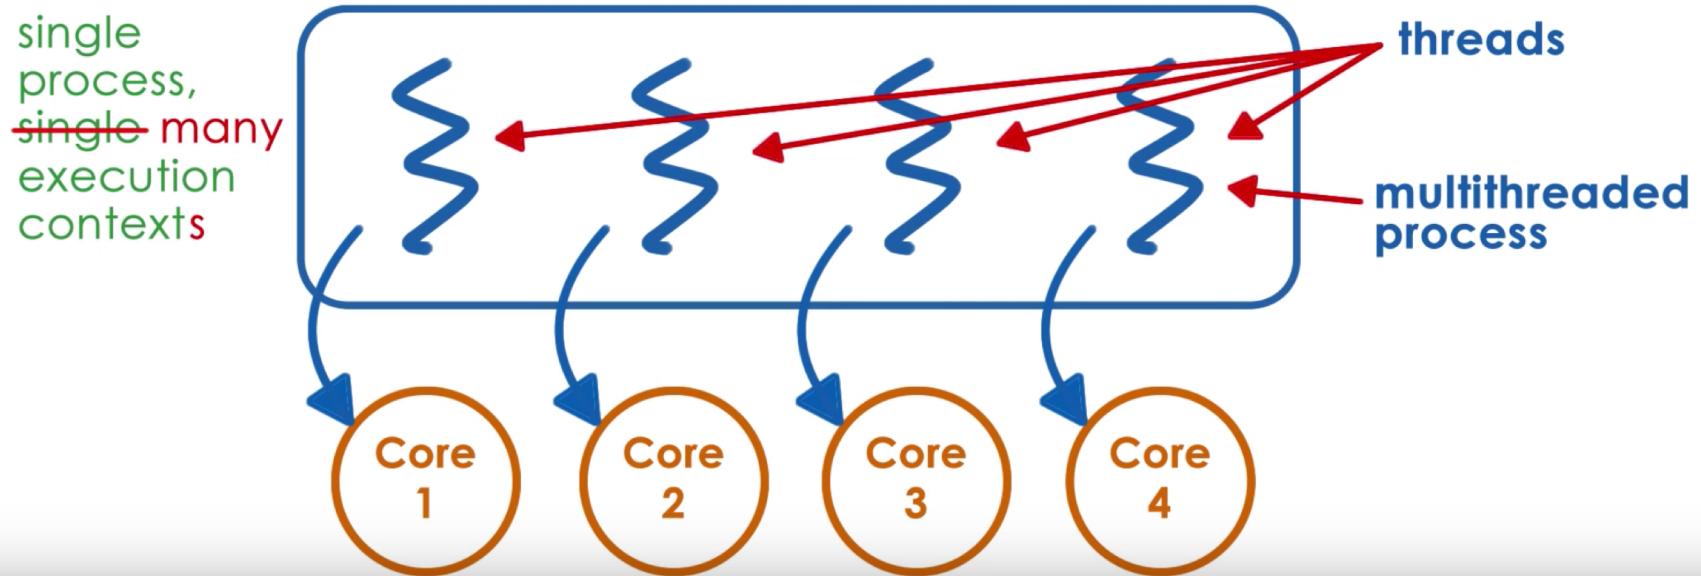


# What if we have multiple CPUs

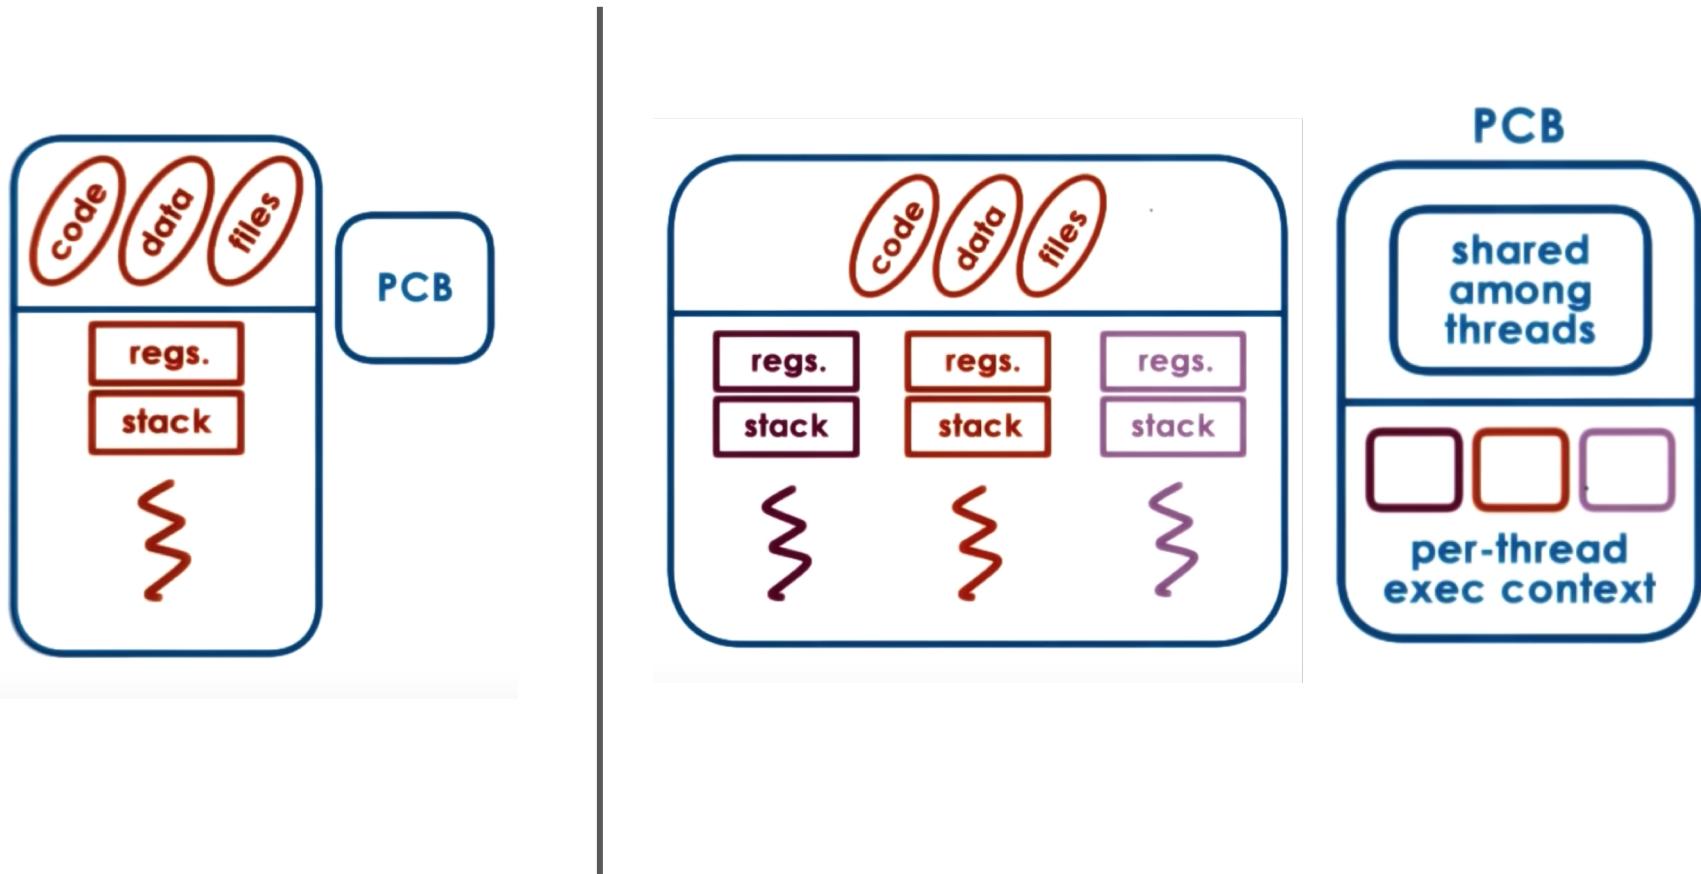
single process,  
single execution  
context



# What if we have multiple CPUs



# Process vs Threads



# Process vs Thread

Process	Thread
Executing instance of an application	Path of execution within a process
Used for heavy weight tasks	Used for small tasks
Takes more time for termination and creation and context switching	Takes less time for termination and creation and context switching
Process consume more resources.	Threads consume less resources
Process is isolated	Threads share memory



# Week 4 – Checklist

- Discuss PA2
- Discuss processes vs threads
- Read about IPC

