



# CSCI-3753: Operating Systems

## Fall 2019

**Abigail Fernandes**

**Department of Computer Science**  
**University of Colorado Boulder**

# Week 6

## Threads

## PThreads

## Shared Memory (IPC)



# Visual Metaphor

A **thread** is like a **worker in a toy shop**!



## Worker in a toy shop

- Is an active entity
  - Executing a unit of toy order
- Works simultaneously with others
  - Many workers completing toy orders
- Requires co-ordination
  - Sharing of tools, parts and workstations



# Visual Metaphor

A **thread** is like a **worker** in a toy shop!

## Thread

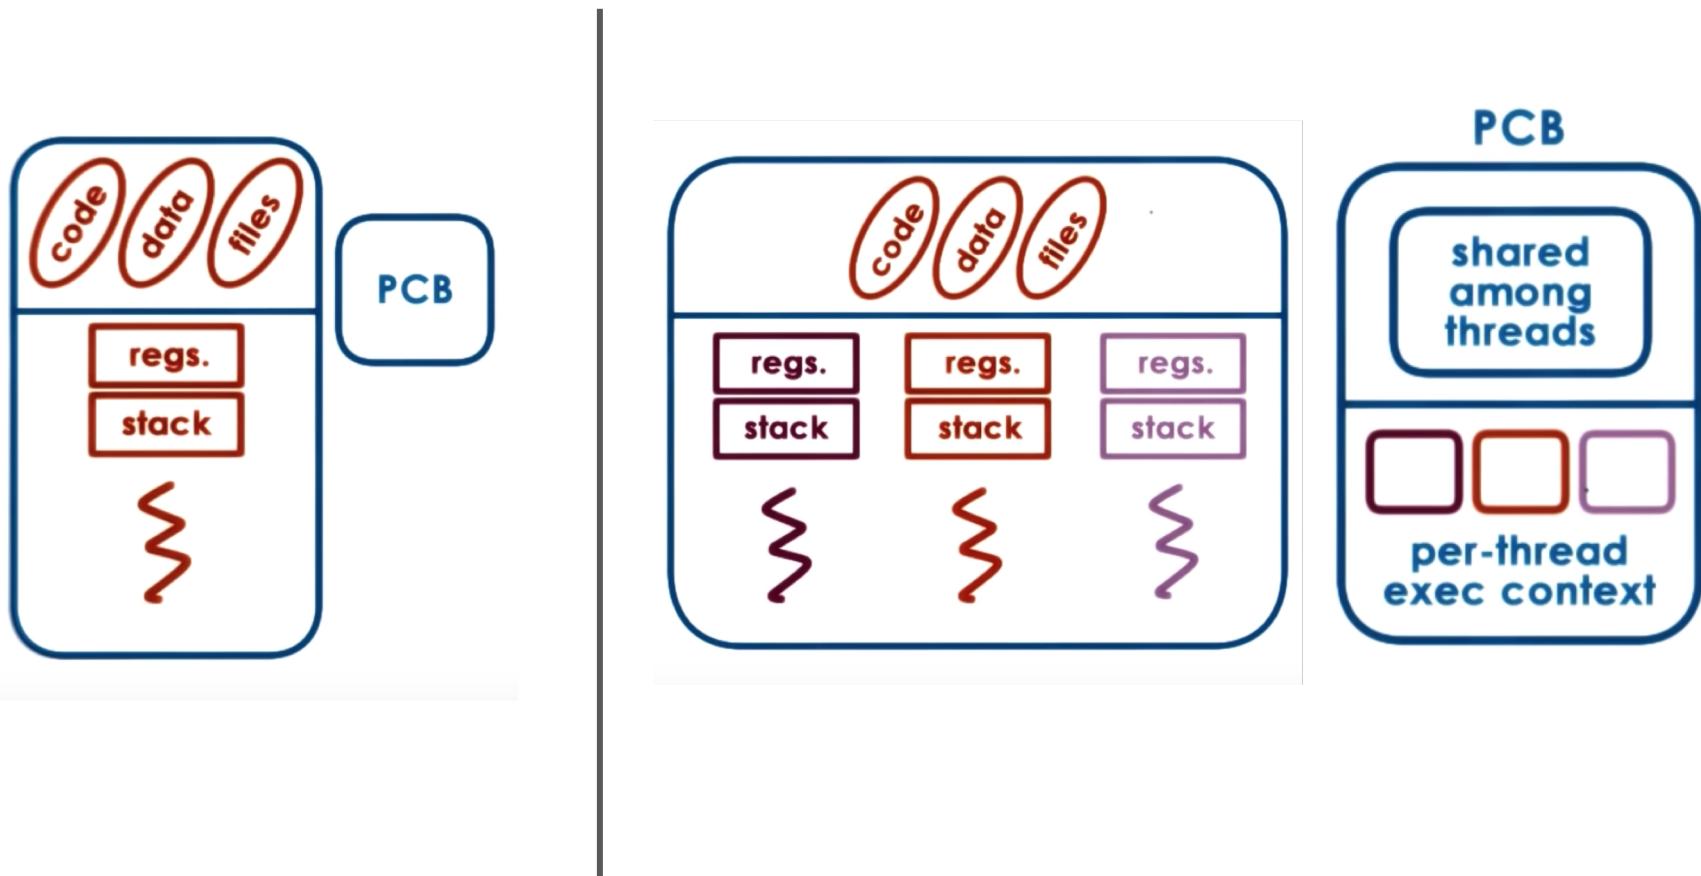
- Is an active entity
  - Executing a unit of process
- Works simultaneously with others
  - Many threads executing
- Requires co-ordination
  - Sharing I/O devices, CPUs, memory

## Worker

- Is an active entity
  - Executing a unit of toy order
- Works simultaneously with others
  - Many workers completing toy orders
- Requires co-ordination
  - Sharing of tools, parts and workstations



# Process vs Threads



# Threads

- Threads **share** the same address space
  - Code
  - Data
  - Other OS resources
    - Heap segments
    - Open file descriptors
    - Signal and signal handlers
    - Current working directory
    - User and group ids
- Each thread has its own **unique**
  - Thread ID
  - Stack
  - Set of registers
  - Program counter



# Common Pitfalls with Multithreading



- Race Conditions
- Thread Safe Code
- Mutex deadlock

# Common Pitfalls with Multithreading

- **Race Conditions**

While the code may appear on the screen in the order you wish the code to execute, threads are scheduled by the operating system and are executed at random

- Thread Safe Code
- Mutex deadlock



# Common Pitfalls with Multithreading

- Race Conditions
- **Thread Safe Code**
  - There should be no static or global variables which other threads may clobber or read assuming single threaded operation.
  - Many non-reentrant functions return a pointer to static data. → Thread-unsafe functions
- Mutex deadlock



# Common Pitfalls with Multithreading

- Race Conditions
- Thread Safe Code
- **Mutex deadlock**
  - Occur when a mutex is applied but then not "unlocked".
  - Cause program execution to halt indefinitely.



# Week 6

## Threads

## PThreads

## Shared Memory (IPC)



# PThreads

- PThreads == **POSIX** Threads
- POSIX = Portable Operating Systems Interface
- POSIX Threads
  - POSIX version of [Birell's API](#)
  - **Specifies** syntax and semantics of the operations
  - API for thread creation and synchronization
- May be provided as user level or kernel level

```
#include <pthread.h>
```



# The PThreads API

- **Thread management:** The first class of functions work directly on threads - creating, terminating, joining, etc.
- **Mutexes:** provide for creating, destroying, locking and unlocking mutexes.
- **Semaphores:** provide for create, destroy, wait, and post on semaphores.
- **Condition variables:** include functions to create, destroy, wait and signal based upon specified variable values.



# Thread Creation

## Birell's Mechanism

- Thread
- Fork (proc, args)
  - Thread creation
- Join (thread)

## PThreads

```
pthread_t aThread; // type of thread
```

```
int pthread_create (thread, attr, start_routine, arg)
```

```
int pthread_join(pthread_t thread, void ** status)
```



# Thread Creation

```
int pthread_create (tid, attr, start_routine, arg)
```

- It returns the new thread ID via the *tid* argument.
- The *attr* parameter is used to set thread attributes, NULL for the default values.
- The *start\_routine* is the C routine that the thread will execute once it is created.
- A single argument may be passed to *start\_routine* via *arg*. It must be passed by reference as a pointer cast of type void.
- Returns status information that indicates if creation was successful or not.



# Thread Termination and Join

```
int pthread_exit (value)
```

```
int pthread_join(pthread_t thread_id,  
void ** status)
```

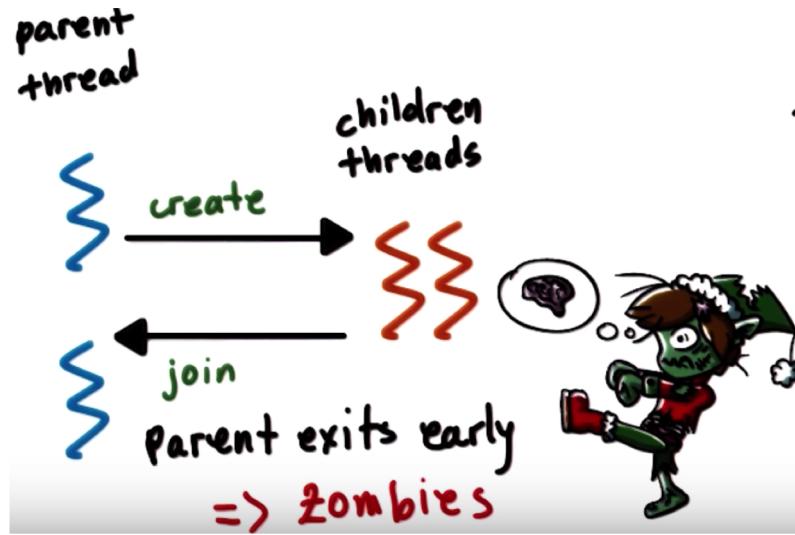
- This function is used by a thread to terminate
- The return value is passed as a pointer.

- This subroutine blocks the calling thread until the specified *thread\_id* thread terminates
- Returns 0 on success, -1 on failure.
- The returned value is a pointer returned by reference.
- If you do not care about the return value, you can pass NULL for the second argument

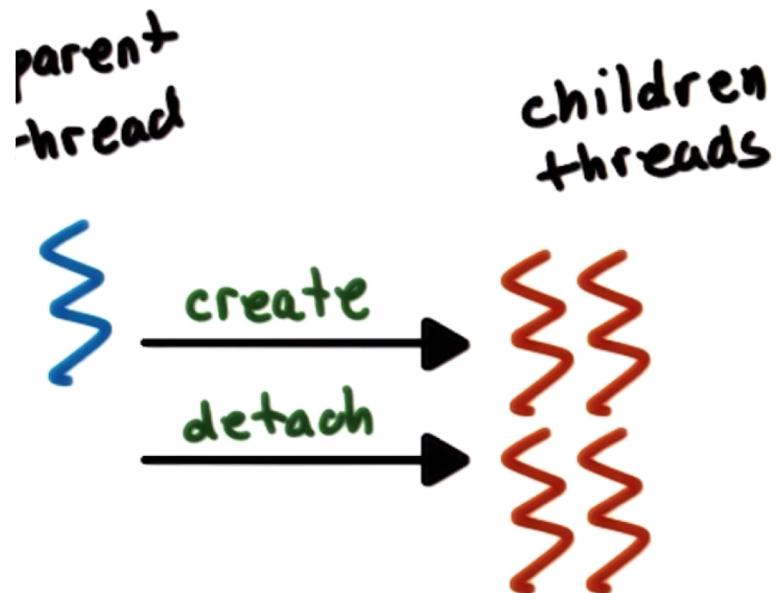


# Detaching PThreads

**Default: Joinable Threads**



**Detached Threads**



# Compiling PThreads

1. `#include <pthread.h>` in main file.

2. Compile source with `-lpthread` or `-pthread`

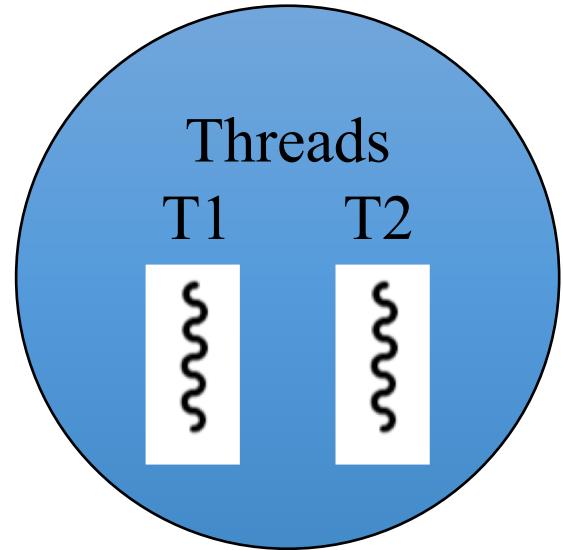
```
Intro to OS ~ ==> gcc -o main main.c -lpthread  
Intro to OS ~ ==> gcc -o main main.c -pthread
```

3. Check return values of common functions



## Example 1

- One main thread creates two new threads with **default** parameters.
- **Main thread:**
  - Write whole numbers in a file called ***whole\_num***
- **New thread 1:**
  - Execute **odd()** function that writes odd numbers in a file called ***odd\_num***
- **New thread 2:**
  - Execute **even()** function that writes even numbers in a file called ***even\_num***

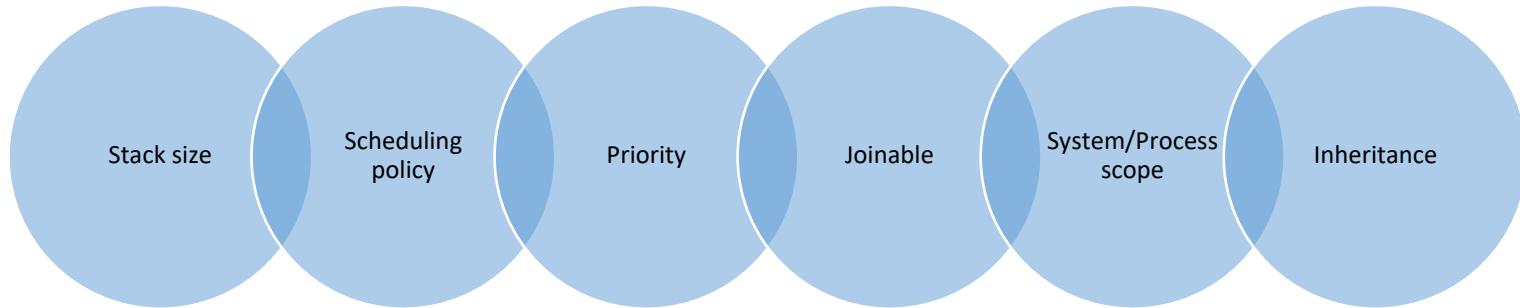


# PThread Attributes

```
int pthread_create(pthread_t *thread,  
                  const pthread_attr_t *attr,  
                  void * (*start_routine)(void *),  
                  void *arg));
```

## pthread\_attr\_t

- Specified in `pthread_create`
- Defines features of the new thread
- Has default behavior with NULL in `pthread_create`



# PThread Attributes

```
int pthread_attr_init(pthread_attr_t *attr);
int pthread_attr_destroy(pthread_attr_t *attr);

int pthread_attr_setstacksize(pthread_attr_t *attr,
size_t sz);
int pthread_attr_getstacksize(const pthread_attr_t
*attr, size_t *sz);

int pthread_attr_setstack(pthread_attr_t *attr, void
*addr, size_t stacksize);
int pthread_attr_getstack(const pthread_attr_t *attr,
void **addr, size_t *stacksize);
```



# Thread Creation (Attributes)

- Thread's SET SCOPE
  - Scope values
    - PTHREAD\_SCOPE\_PROCESS (unbound, default)
      - The thread competes for resources with all other threads in the same process.
      - They are scheduled relative to one another according to their scheduling policy and priority.
    - PTHREAD\_SCOPE\_SYSTEM (bound)

```
int pthread_attr_setscope(pthread_attr_t *attr, int scope);
int pthread_attr_getscope(const pthread_attr_t *attr, int *scope);
```



# Thread Creation (Attributes)

- Thread scheduling attributes
  - Inherit values
    - PTHREAD\_INHERIT\_SCHED (default)
      - Threads that are created using attr inherit scheduling attributes from the creating thread.
    - PTHREAD\_EXPLICIT\_SCHED
      - Threads that are created using attr take their scheduling attributes from the values specified by the attributes object.

```
int pthread_attr_setinheritsched(pthread_attr_t *attr,  
int inherit);  
int pthread_attr_getinheritsched(const pthread_attr_t  
*attr, int *inherit);
```



# Thread Creation (Attributes)

- Thread scheduling attributes
  - Policy values
    - SCHED\_FIFO (First in-first out scheduling)
    - SCHED\_RR (Round-robin scheduling)
    - SCHED\_OTHER (Default Linux time-sharing scheduling)

```
int pthread_attr_setschedpolicy(pthread_attr_t  
*attr, int policy);  
int pthread_attr_getschedpolicy(const pthread_attr_t  
*attr, int *policy);
```



## Process P1

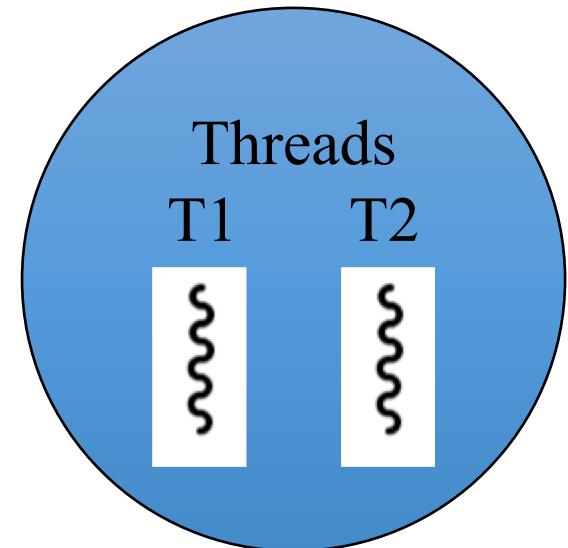
### Example 2

One main thread creates two new threads with **specific** parameters.

All 3 threads write to standard output  
(i.e., only monitor).

All 3 threads use

- Case 1: Local variables -> predict the execution
- Case 2: The same global variable -> predict the execution



# PThread Mutexes

Solve mutual exclusion problems among concurrent threads.

## Birell's Mechanism

- Mutex
- Lock (mutex) {  
}

## PThreads

```
pthread_mutex_t aMutex; // mutex type
```

```
int pthread_mutex_lock (pthread_mutex_t *mutex)
```

```
int pthread_mutex_unlock (pthread_mutex_t *mutex)
```



# Pthread Mutexes

Birell

```
list<int> my_list;  
Mutex m;  
  
void safe_insert(int i) {  
    Lock(m) {  
        my_list.insert(i);  
    } // unlock;  
}
```

PThreads

```
list<int> my_list;  
pthread_mutex_t m;  
  
void safe_insert(int i) {  
    pthread_mutex_lock(m);  
    my_list.insert(i);  
    pthread_mutex_unlock(m);  
}
```



# Other Mutex Operations

```
int pthread_mutex_init(pthread_mutex_t *mutex,  
                      const pthread_mutexattr_t *attr);  
// mutex attributes == specifies mutex behavior when  
// a mutex is shared among processes
```

```
int pthread_mutex_trylock(pthread_mutex_t *mutex);
```

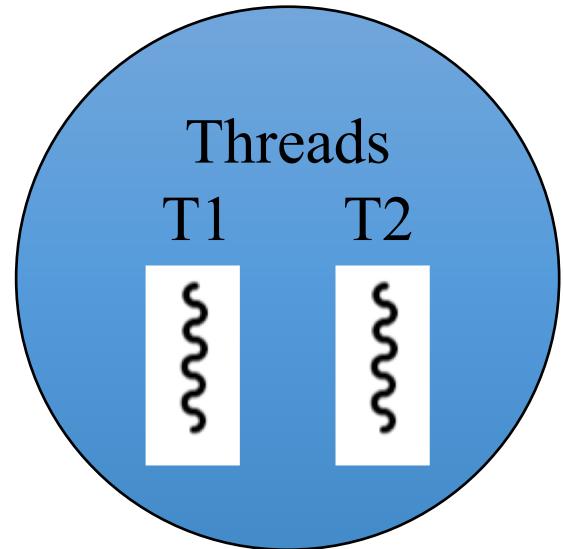
```
int pthread_mutex_destroy(pthread_mutex_t *mutex);
```

A null value of *attr* initializes mutex with default attributes.



## Example 3 (Mutex)

- One main thread creates two new threads with **specific** parameters.
- All 3 threads write to standard output (i.e., only monitor).
- The threads use **mutexes** to synchronize their writes in the following 2 cases:
  - Case 1:** Local variables -> predict the execution
  - Case 2:** The same global variable -> predict the execution



# Week 6

## Threads

## PThreads

## Shared Memory (IPC)



# Visual Metaphor

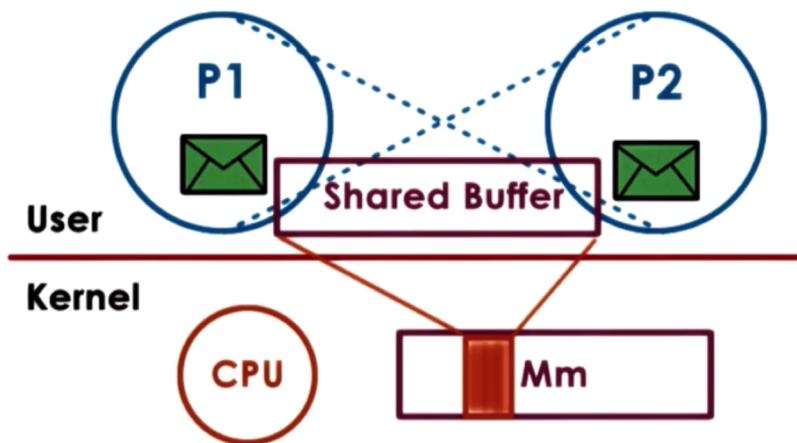
IPC is like working together in a toy shop



- **Workers share work area**
  - Parts and tools on the table
- **Workers call each other**
  - Explicit requests and responses
- **Requires synchronization**
  - I'll start when you finish



# Shared Memory IPC



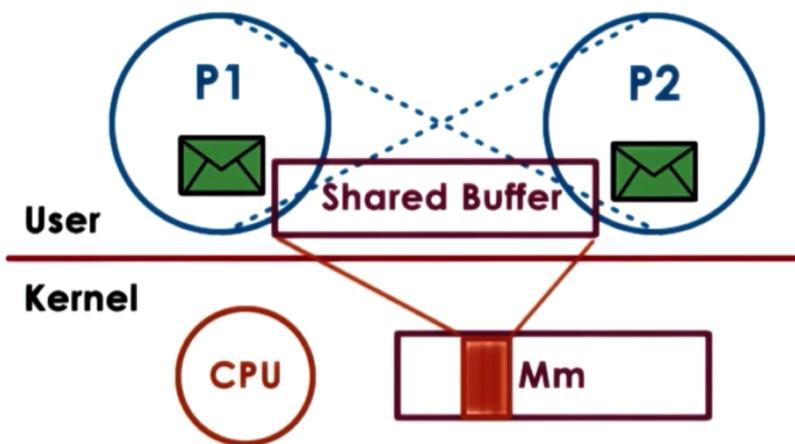
Read and write to shared memory region

OS establishes shared channel between the processes.

1. Physical pages mapped into virtual address space
2.  $VA(p1)$  and  $VA(p2)$  map to the same physical address.
3.  $VA(p1) \neq VA(p2)$
4. Physical memory does not need to be contiguous



# Shared Memory IPC



- Explicit synchronization
- Programmers responsibility: Communication protocol, shared buffer management
- Complicated to implement in a distributed system



- System calls only for set up, so very fast.
- Data copies are potentially reduced (but not eliminated)
- Useful for large amounts of data



# Shared Memory

- Create a shared memory segment
- `int shmget(key_t key, size_t size, int shmflg);`
- Other processes (if permitted) can access.
- Library: `sys/ ipc.h` and `sys/shm.h`
  - `key`: a value used to share between processes
  - `size`: size of the shared memory segment
  - `shmflg`: options for segment creation



# Shared Memory

- Attach the shared segment to a process address space

```
void *shmat(int shmid, const void *shmaddr,...int  
shmflg);
```

- Once attached, a process can read or write to the segment.
- Library: *sys/types.h* and *sys/shm.h*
  - *shmid*: id of the shared memory segment
  - *shmaddr*: address in the process space
  - *shmflg*: options for shared memory permission



# Shared Memory

- Detach the shared memory segment located at the `shmaddr` address

```
int shmdt(const void *shmaddr);
```

- Library: `sys/types.h` and `sys/shm.h`
  - `shmaddr`: attaching address in the process space

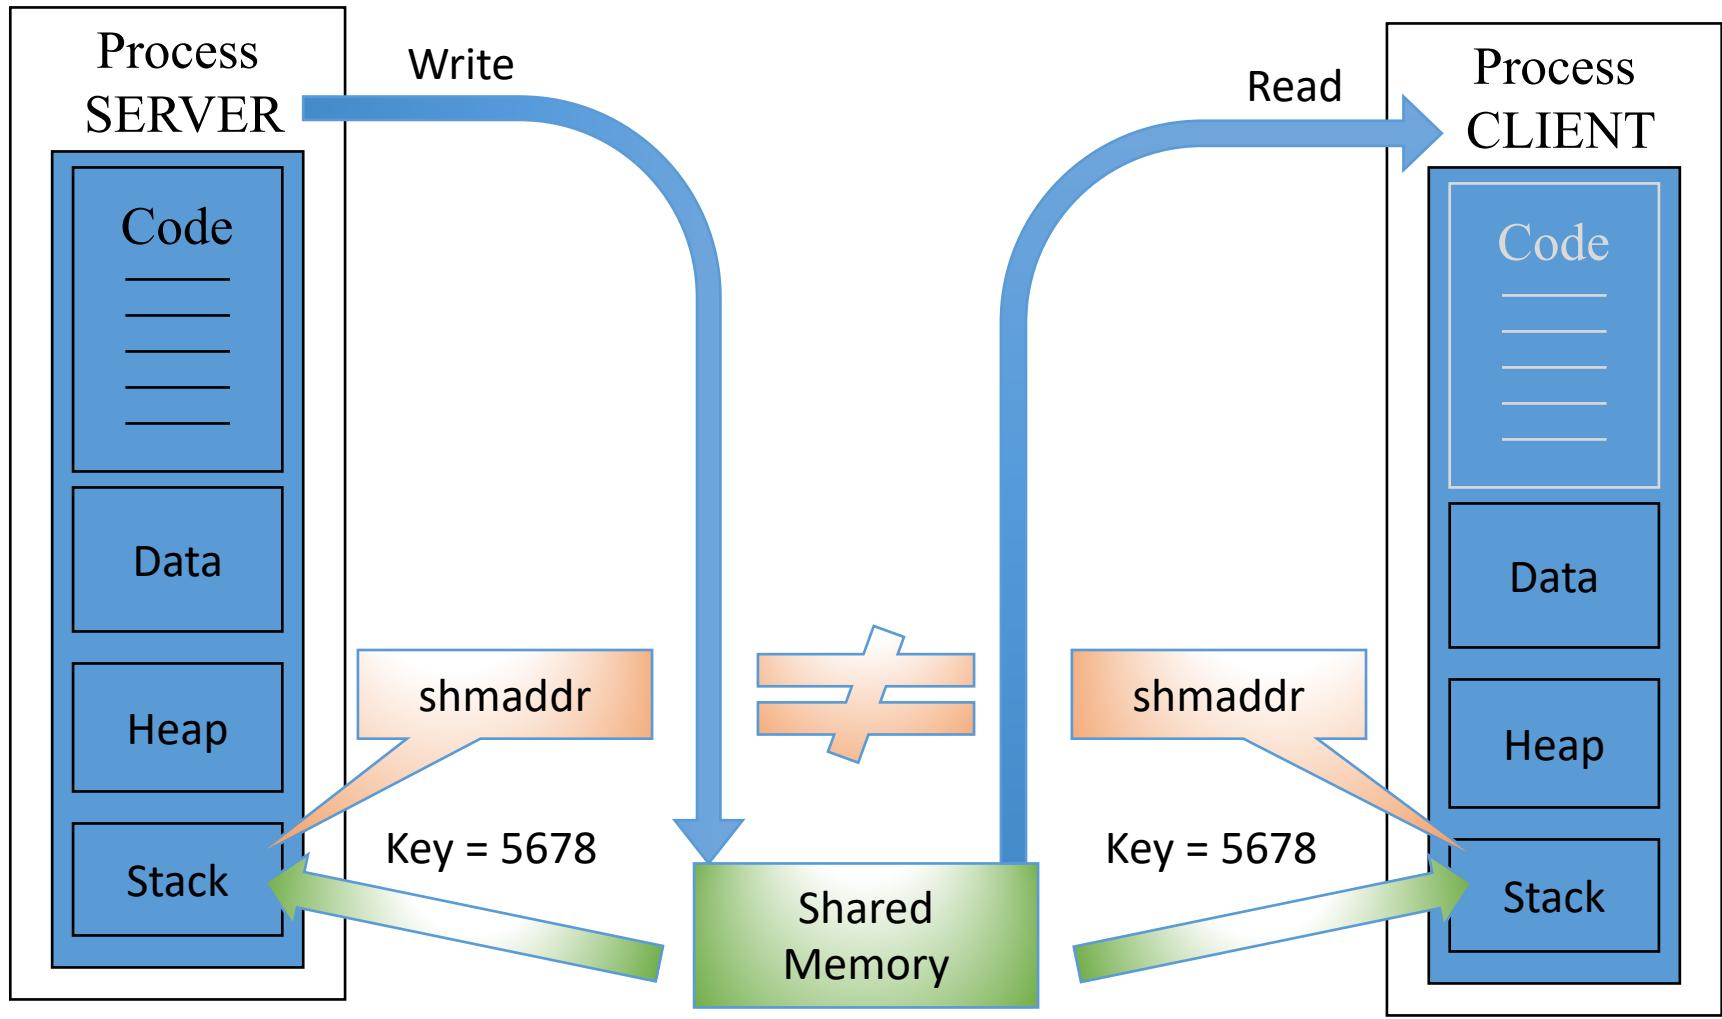


# Shared Memory Example

- Goal: To illustrate the usage of shared memory between two processes in Linux.
- The **server** process creates a shared memory segment and writes some characters in it. It then waits for the first character in the shared memory to become '\*', and then exits.
- The **client** process reads characters from this shared memory and prints them out on the terminal, writes '\*' at the beginning of the shared memory segment, and then exits.



# Shared Memory Example (Code Demo)



# Week 6 – Checklist

- Discuss PThreads
- Reviewed examples on PThreads
- Discuss Shared Memory from IPC
- Read more about threads
- BEGIN PA3!!!!**

