



CSCI-3753: Operating Systems

Fall 2019

Abigail Fernandes

Department of Computer Science
University of Colorado Boulder

Week 11

- > Programming Assignment 4
- > Memory Management



Assignment Goal

Implement a paging strategy that a paging simulator can use to maximize the performance of the memory access in a set of pre-defined programs

Action items

- Implement **LRU** algorithm
[`pager-lru.c`](#)
- Implement any form of **predictive paging algorithm**
[`pager-predict.c`](#)



Paging Simulator

- Run a random set of 5 pre-defined programs utilizing a limited number of shared physical pages
- Provided default values in `simulator.h`
 - 20 virtual pages per process (`MAX_PROC_PAGES`)
 - 100 physical pages (frames) in total (`PHYSICAL_PAGES`)
 - 20 `simultaneous` processes competing for pages (`MAX_PROCESSES`)
 - 128 memory unit page size (`PAGE_SIZE`)
 - `100 tick delay` to swap a page in or out (`PAGE_WAIT`)
 - Each instruction or step in the simulated programs requires 1 tick to complete.



Paging Simulator

- Is the environment resource constrained?
- How many physical pages can be swapped in at a given time?
- How many virtual pages will you have to access at most?
- Swapping a page in and out takes how much time?



Paging Simulator

Provide 3 functions for interaction

- To control the allocation of virtual and physical pages
 - `pagein()`
 - `pageout()`
- To handle the page fault
 - **`pageit()`** ← core paging function that needs implementation

Source code is provided.

- `simulator.h, simulator.c`
- `pager-basic.c, pager-lru.c, pager-predict.c`

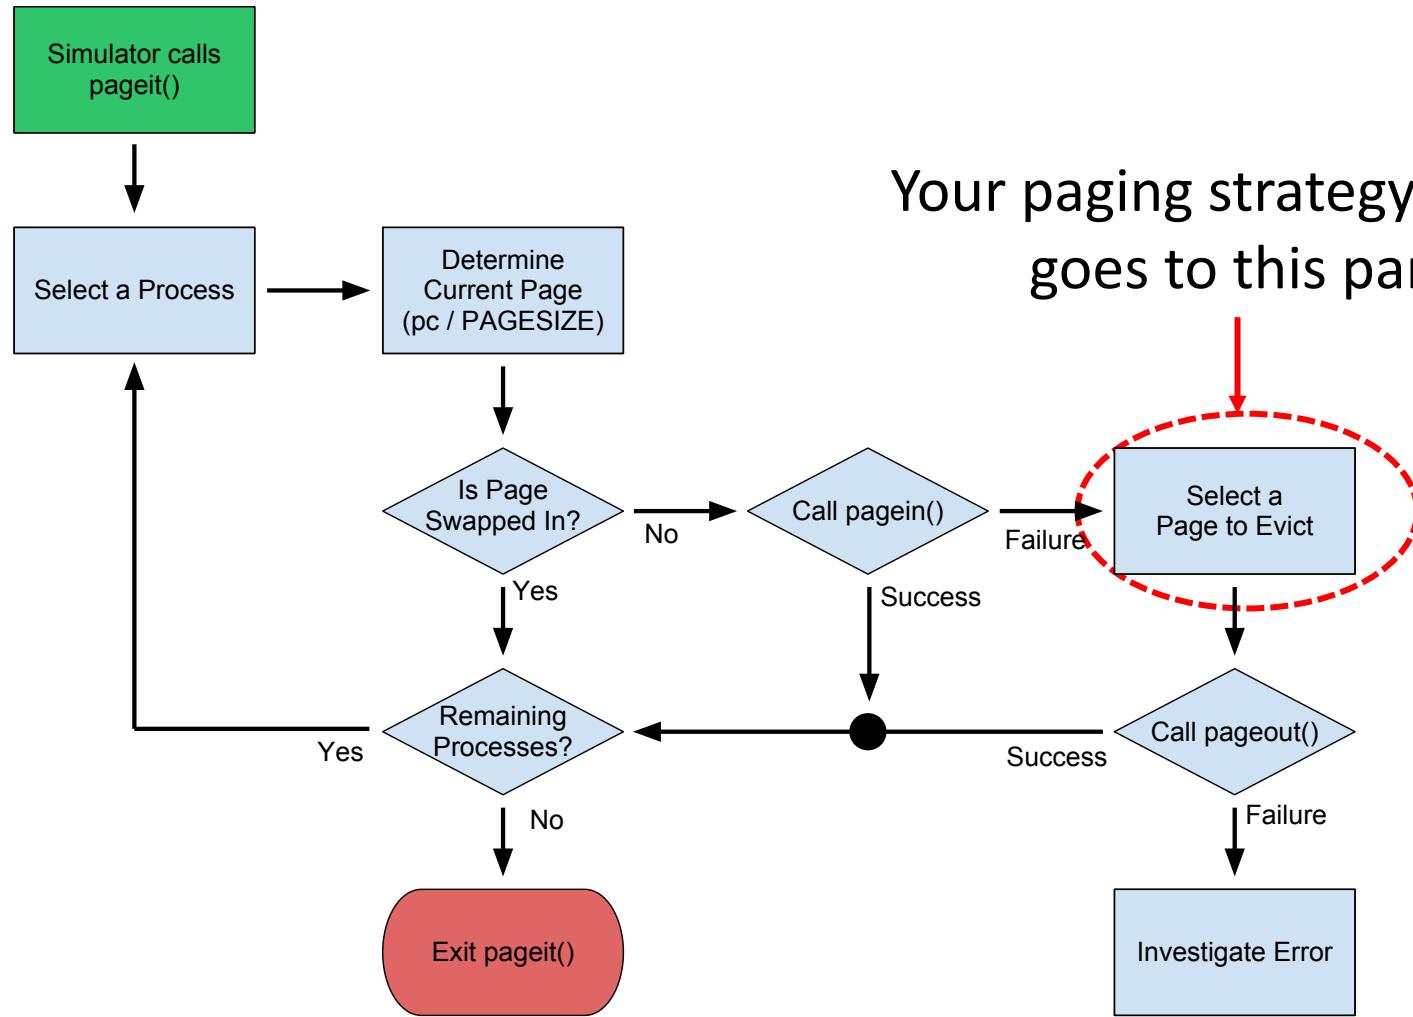


pager-basic.c

- A basic “one-process-at-a-time” implementation
- A simple demonstration of the simulator API
- Doesn’t need any implementation from YOU!!!



pager-basic.c



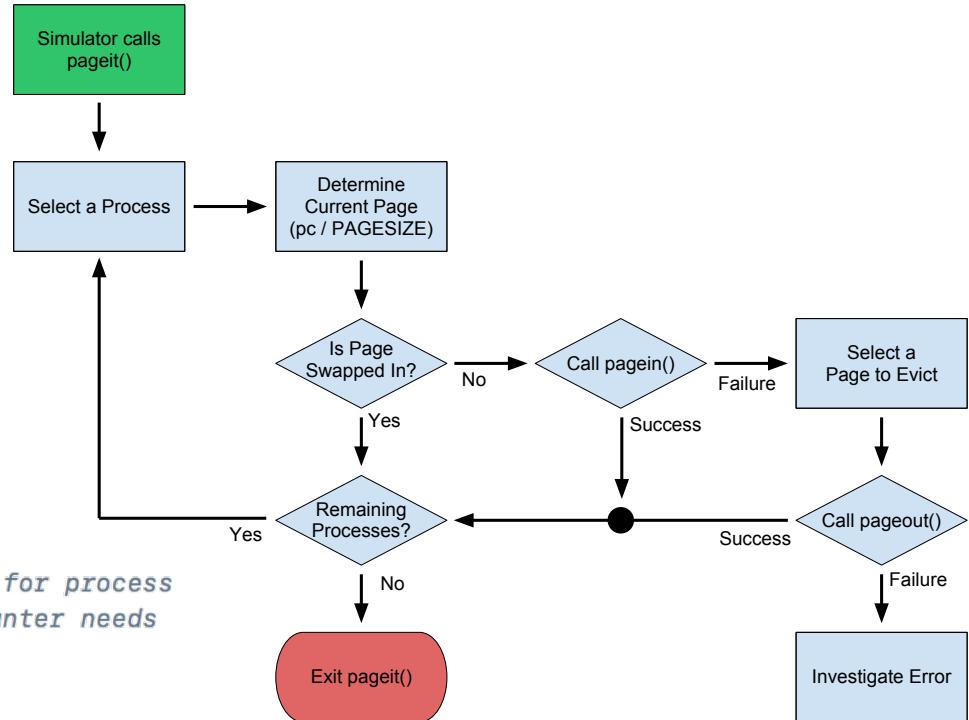
```

#include "simulator.h"

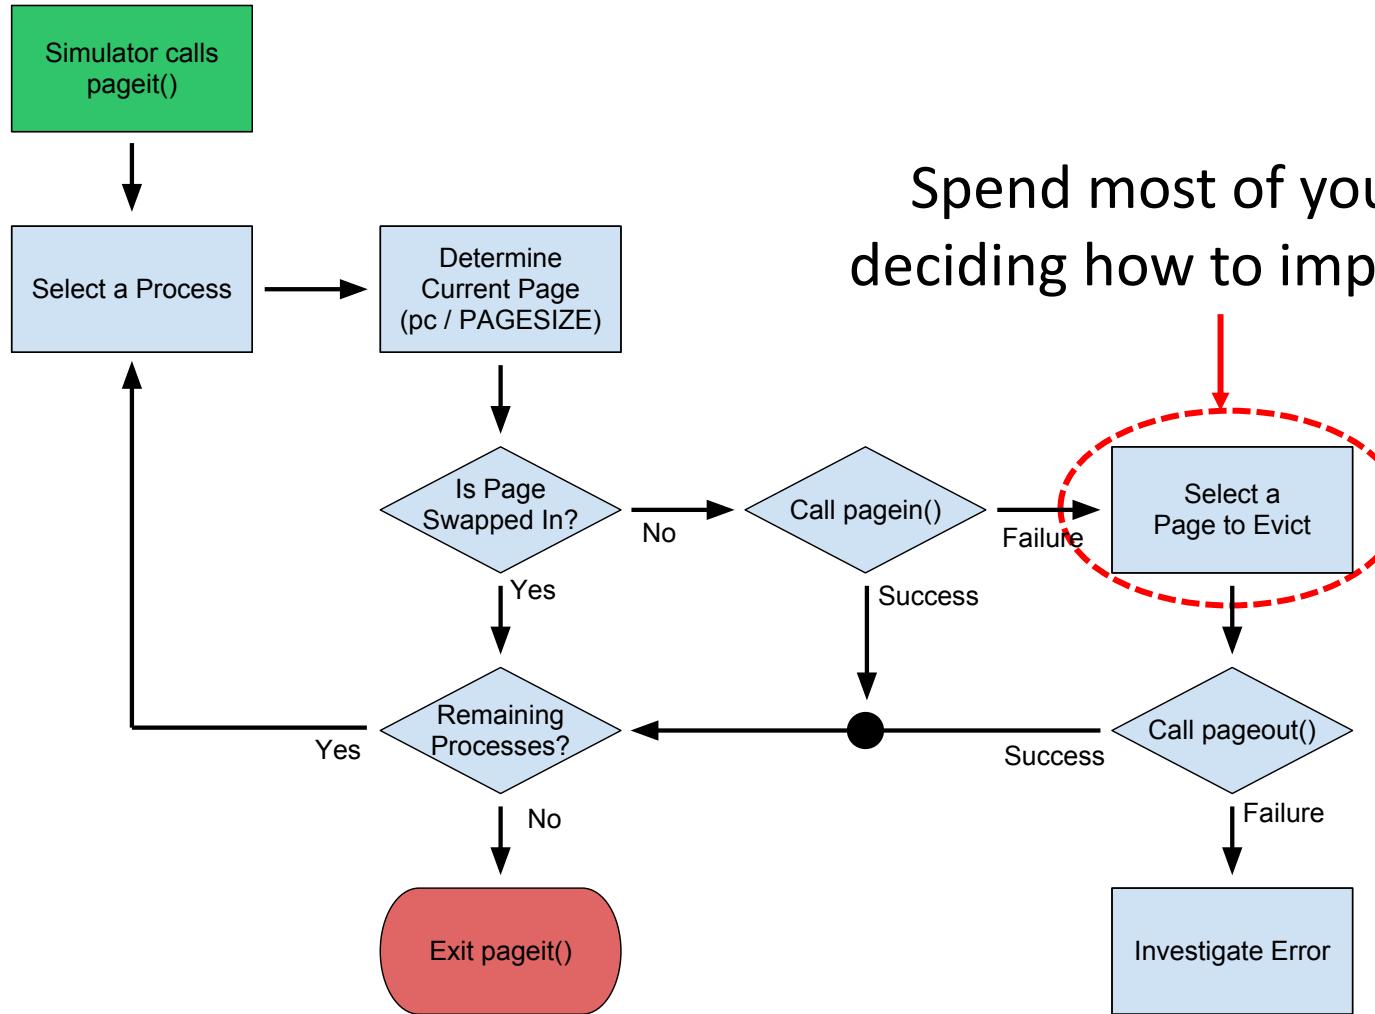
void pageit(Pentry q[MAXPROCESSES]) {
    /* Local vars */
    int proc;
    int pc;
    int page;
    int oldpage;

    /* Trivial paging strategy */
    /* Select first active process */
    for(proc=0; proc<MAXPROCESSES; proc++) {
        /* Is process active? */
        if(q[proc].active) {
            /* Dedicate all work to first active process*/
            pc = q[proc].pc;           // program counter for process
            page = pc/PAGESIZE;       // page the program counter needs
            /* Is page swaped-out? */
            if(!q[proc].pages[page]) {
                /* Try to swap in */
                if(!pagein(proc,page)) {
                    /* If swapping fails, swap out another page */
                    for(oldpage=0; oldpage < q[proc].npages; oldpage++) {
                        /* Make sure page isn't one I want */
                        if(oldpage != page) {
                            /* Try to swap-out */
                            if(pageout(proc,oldpage)) {
                                /* Break loop once swap-out starts*/
                                break;
                            }
                        }
                    }
                }
            }
            /* Break loop after finding first active process */
            break;
        }
    }
}

```



pager-lru.c



Spend most of your time
deciding how to implement it



```

#include <stdio.h>
#include <stdlib.h>

#include "simulator.h"

void pageit(Pentry q[MAXPROCESSES]) {

    /* This file contains the stub for an LRU pager */
    /* You may need to add/remove/modify any part of this file */

    /* Static vars */
    static int initialized = 0;
    static int tick = 1; // artificial time
    static int timestamps[MAXPROCESSES][MAXPROCPAGES];

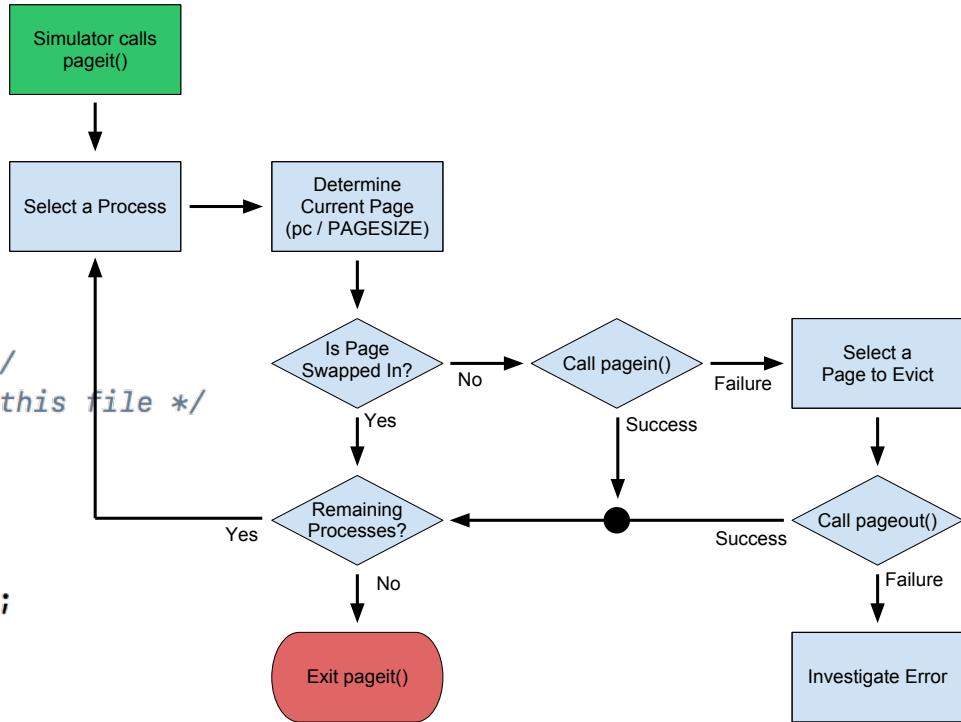
    /* Local vars */
    int proctmp;
    int pagetmp;

    /* initialize static vars on first run */
    if(!initialized){
        for(proctmp=0; proctmp < MAXPROCESSES; proctmp++){
            for(pagetmp=0; pagetmp < MAXPROCUPAGES; pagetmp++){
                timestamps[proctmp][pagetmp] = 0;
            }
        }
        initialized = 1;
    }

    /* TODO: Implement LRU Paging */
    fprintf(stderr, "pager-lru not yet implemented. Exiting...\n");
    exit(EXIT_FAILURE);

    /* advance time for next pageit iteration */
}

```



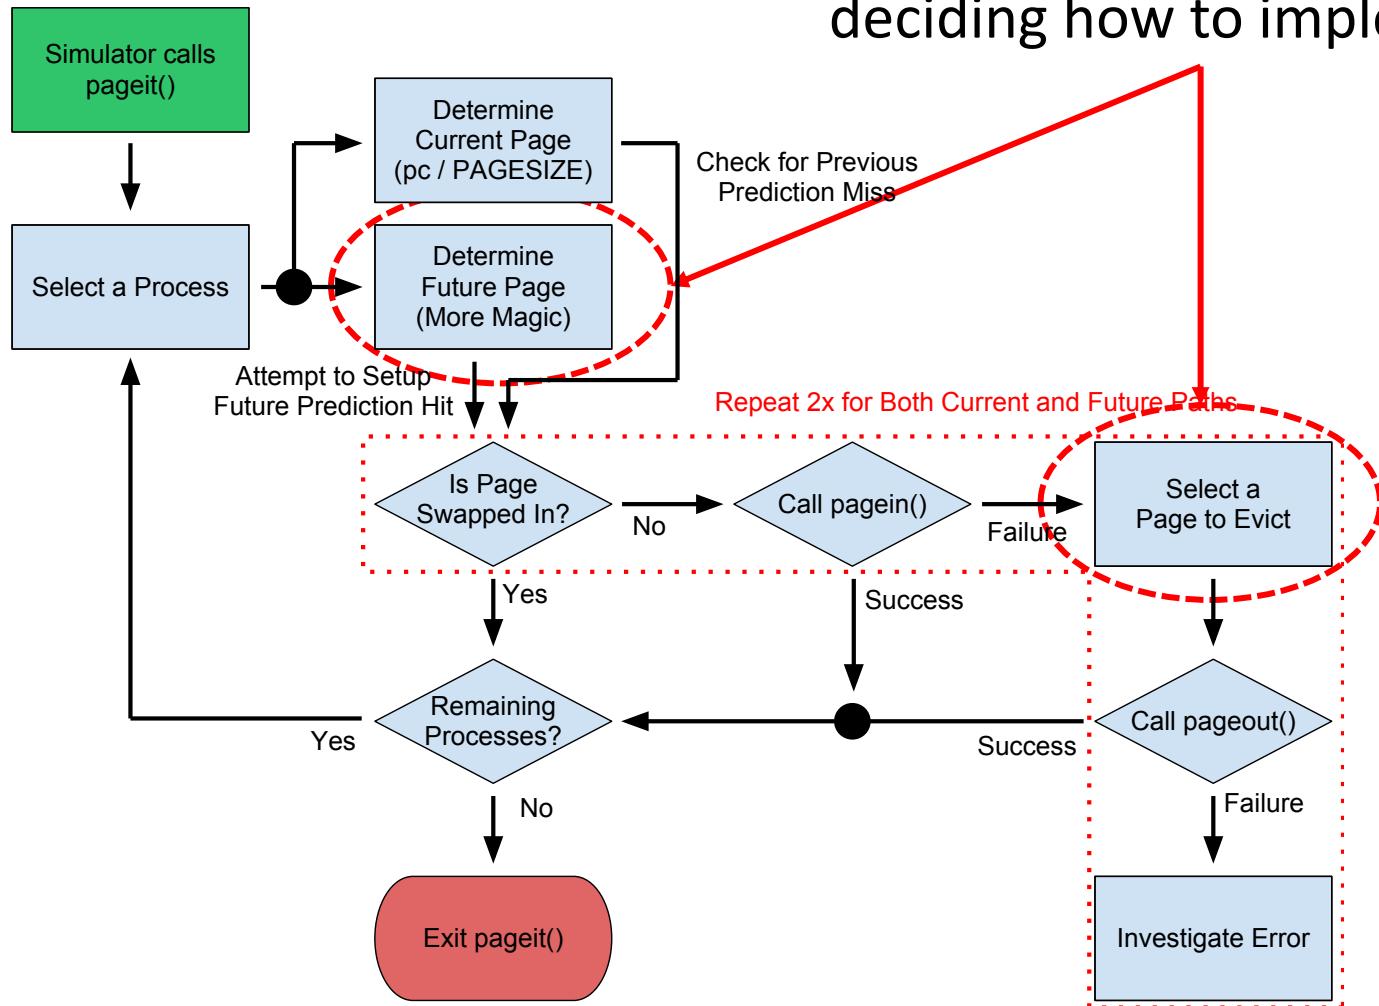
pager-predict.c

- Require a predictive algorithm that
 - Attempts to predict what pages each process will require in the future and then
 - Swaps these pages in before they are needed
- **Note:** In any predictive operation, you ideally wish to stay 100-200 ticks ahead of the execution of each process.



pager-predict.c

Spend most of your time
deciding how to implement it



```

#include <stdio.h>
#include <stdlib.h>

#include "simulator.h"

void pageit(Pentry q[MAXPROCESSES]) {
    /* This file contains the stub for a predictive pager */
    /* You may need to add/remove/modify any part of this file */

    /* Static vars */
    static int initialized = 0;
    static int tick = 1; // artificial time

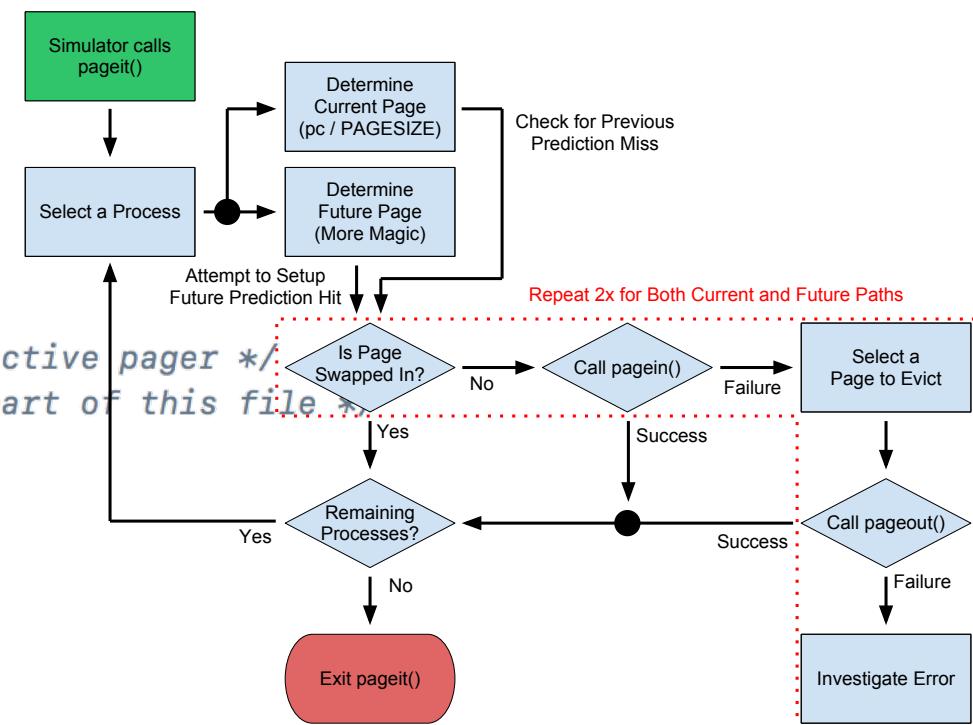
    /* Local vars */
    /* initialize static vars on first run */
    if(!initialized){
        /* Init complex static vars here */

        initialized = 1;
    }

    /* TODO: Implement Predictive Paging */
    fprintf(stderr, "pager-predict not yet implemented. Exiting...\n");
    exit(EXIT_FAILURE);

    /* advance time for next pageit iteration */
    tick++;
}

```



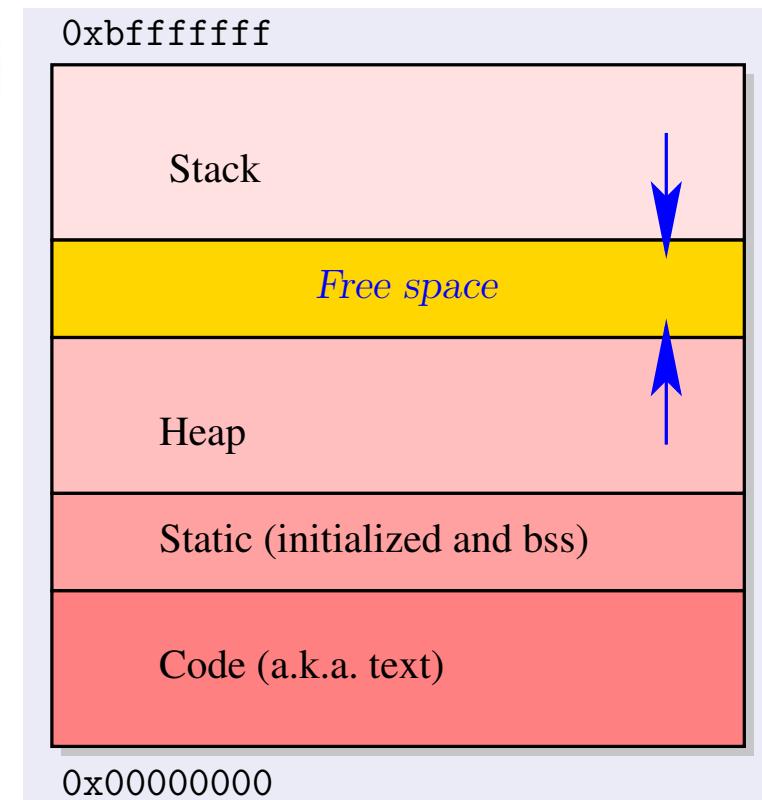
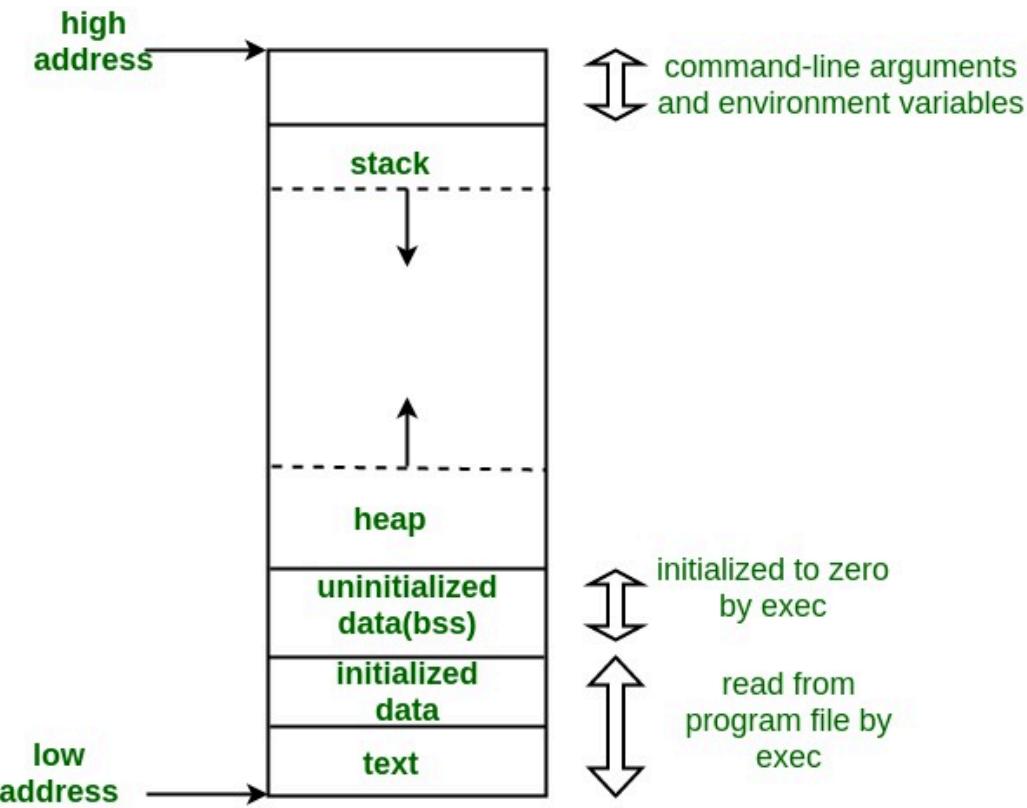
Week 12

- > Programming Assignment 4
- > Memory Management



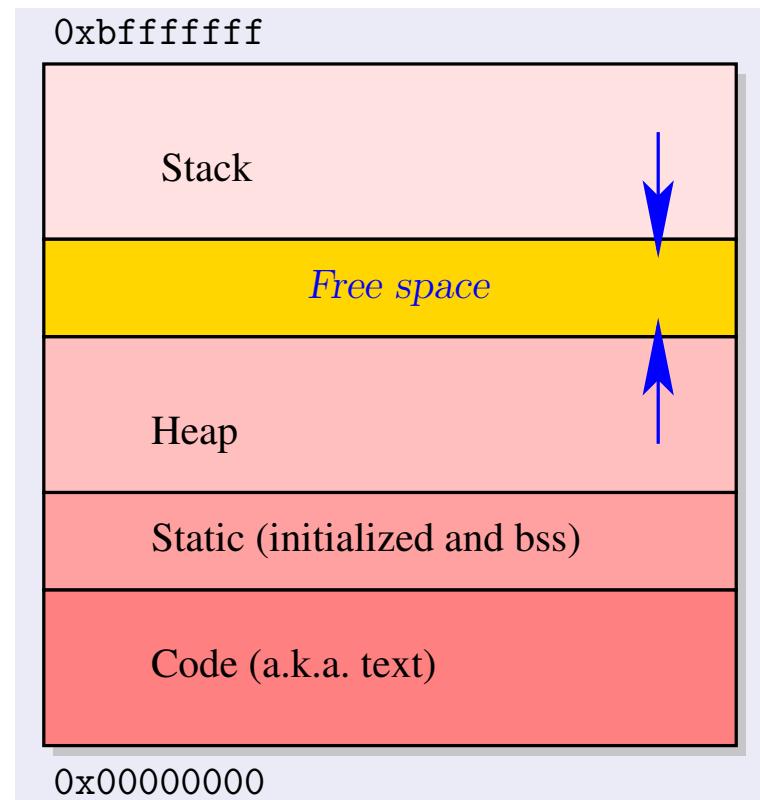
Memory Address Space of a Process

- Allocated and initialized when loading and executing



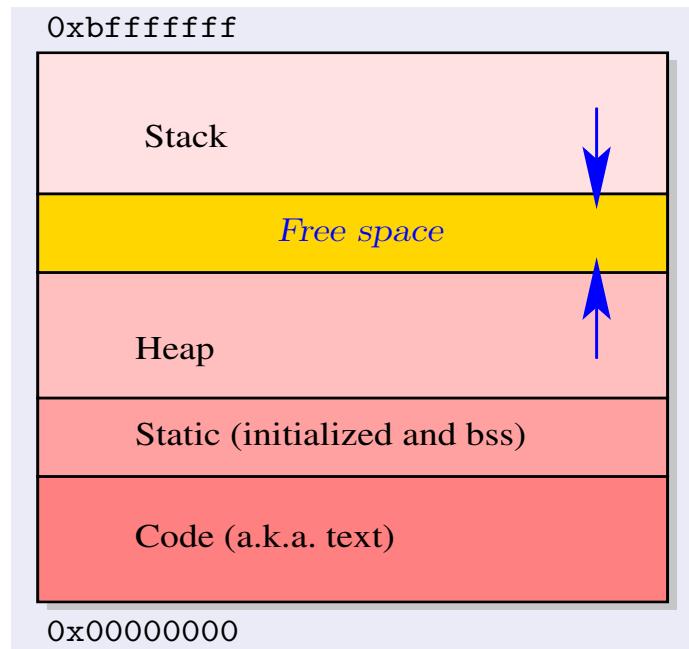
Memory Address Space of a Process

- Allocated and initialized when loading and executing a program
- Consists of
 - Code segment**
 - .o and executable code
 - Below the heap or stack in order to prevent heaps and stack overflows from overwriting it
 - Usually
 - Sharable
 - Read-only
 - Fixed size



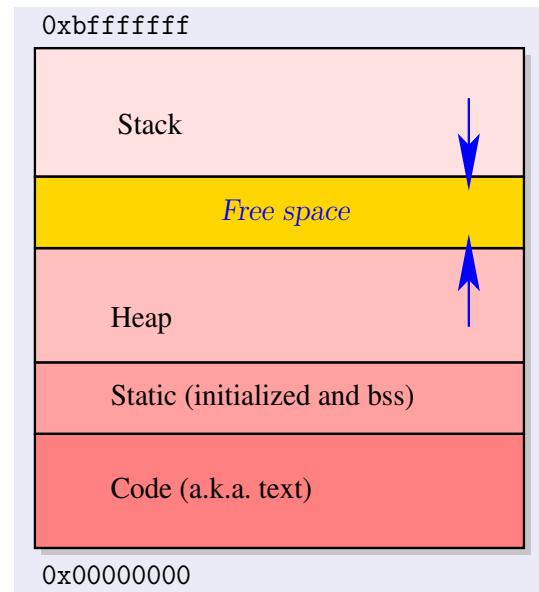
Memory Address Space of a Process

- Allocated and initialized when loading and executing a program
- Consists of
 - Static segments**
 - Data segment**
Initialized global and C static Variables
 - BSS segment**
Uninitialized global variables that are zeroed when initializing the process, are stored in .bss segment



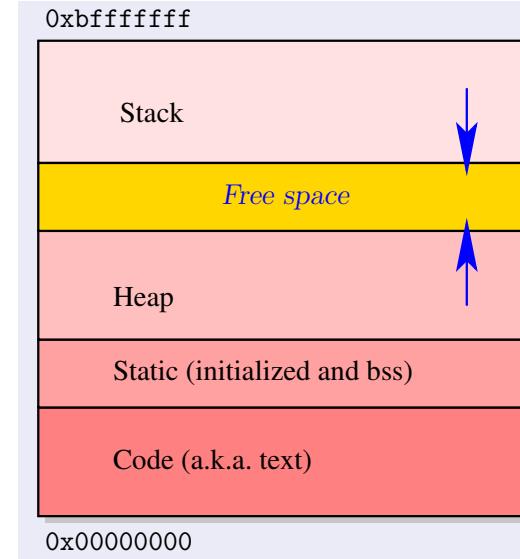
Memory Address Space of a Process

- Allocated and initialized when loading and executing a program
- Consist of
 - Stack segment**
 - Stack frames of function call arguments and local variables, also called automatic variables in C
 - Heap segment**
 - Dynamic allocation
 - malloc, realloc, and free,
 - brk and sbrk system calls to adjust its size

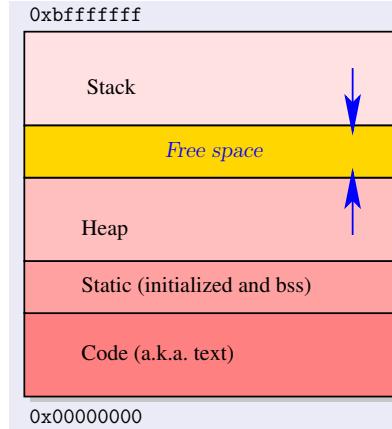


System Calls

- To change **data segment** file
 - `brk()`
 - `sbrk()`
- **Program break** is the first location after the end of the uninitialized bss segment.
 - Increasing the program break has the effect of allocating memory to the process
 - Decreasing the break deallocates memory.
- Both functions are **deprecated as “programmer interface” functions.**

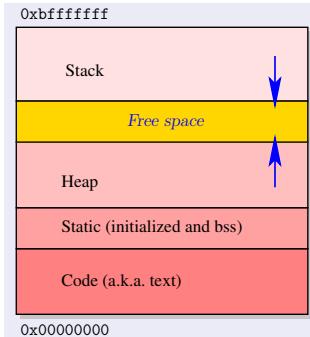


System Calls



```
int brk(void *addr);
```

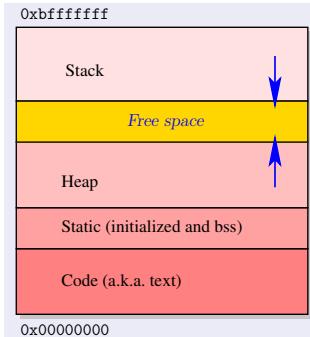
Set the end of the data segment to the value specified by `addr`, when that value is reasonable, the system has enough memory and the process does not exceed its maximum data size



System Calls

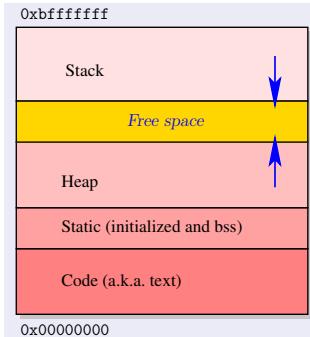
```
void *sbrk(intptr_t increment);
```

- `sbrk()` increments the program's data space by `increment` bytes.
- `sbrk(0)` returns the current location of the program break.



System Calls

- Both `brk()` and `sbrk()` extend heap ???
 - `brk(b)` sets the end of the heap to b
 - `sbrk(n)` extends the end of the heap by n bytes
 - `sbrk(0)` returns the virtual address just past the end of the heap.



System Calls

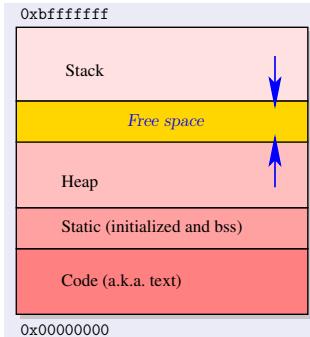
```
void* malloc(size_t size);
```

- Allocate `size` bytes of uninitialized storage
 - Call `sbrk()` to get the memory to allocate in the heap
- More efficient than allocating memory using simply `brk()` or `sbrk()`
 - Do buffering
 - However, does not always invoke `sbrk()`
 - When it calls `sbrk()`, it calls it to allocate a much larger memory than needed

Virtual Memory

- `getpagesize()`
- `getrlimit()` and `setrlimit()`
 - Process virtual memory size limit
 - Max CPU time
 - Max data segment size
 - ...
- `pmap pid/pids`
 - report memory map of a process or processes
- `getconf`
 - get configuration values of standards set in the current operating environment





System Calls

```
void * mmap(void *start, size_t length, int  
prot, int flags, int fd, off_t offset)
```

Create a new mapping in the virtual address space of the calling process

Week 11 – Checklist

- Discuss PA4
- Discuss memory management

